

Contents

1	Introduction	1
1.1	Lean first steps	3
2	Logic and Proposition as Types	4
2.1	First Order Logic	4
2.2	Primitive Types	6
2.3	Curry Howard isomorphism	10
2.4	Predicate logic and dependency	12
2.5	Constructive Mathematics	15
3	Describing and use properties	16
3.1	Exploring Mathlib (The Rat structure)	18
3.2	Coercions and Type Casting	26
3.3	Type Classes and Algebraic Hierarchy in Lean	26
4	Formalizing the topologist's sine curve	30
4.1	T is connected	31
4.1.1	Limits and Convergence with Filters	36
4.1.2	Finalising the first part of the proof	39
4.2	T is not path-connected	39
4.3	T is connected not path-connected	45

1 Introduction

This serves as a brief starting point for understanding how mathematical proofs can be formalized in Lean, as well as being an introduction to the language itself. Lean is both a **functional programming language** and a **theorem prover**. We'll focus primarily on its role as a theorem prover. But what does this mean, and how can that be achieved?

A programming language defines a **set of rules, semantics, and syntax** for writing programs. To achieve a goal, a programmer must write a program that meets given specifications. There are two primary approaches: **program derivation** and **program verification** ([NPS90] Section 1.1). In **program verification**, the programmer first writes a program and then proves it meets the specifications. This approach checks for errors at **run-time** when the code executes. In **program derivation**, the programmer writes a proof that a program with certain properties exists, then extracts a program from that proof. This approach enables specification checking at **compilation-time**, catching errors while typing, thus, before execution. This distinction corresponds to **dynamic** versus **static type systems**. Most programming languages combine both approaches; providing basic types for annotation and compile-time checking, while leaving the remaining checks to be performed at runtime.

Example 1.1 *In a dynamically typed language, like JavaScript, variables can change type after they are created. For example, a variable defined as a number can later be reassigned to a string: TypeScript is a statically typed superset of JavaScript. Unlike JavaScript, it performs type checking at compile time. This means we can prevent the previous behavior, while writing our code, simply by adding a type annotation: Now converting a variable with type annotation number to a string results into a compile time error.*

Nonetheless, TypeScript, even though it has a sophisticated type system, cannot fully capture complex mathematical properties. As well as for the most programming languages, program specifications can only be enforced at runtime. Lean, by contrast, uses a much more powerful type system that enables it to express and verify mathematical statements with complete rigor fully during compilation time. This makes it particularly suitable as a **theorem prover** for formalizing mathematics.

Lean’s type system is based on **dependent type theory**, specifically the **Calculus of Inductive Constructions** (CIC) with various extensions. It’s important to note that **type theory** is not a single, unified theory, but rather a family of related theories with various extensions, ongoing developments, and rich historical ramifications.

Type theory emerged as an alternative foundation for mathematics, addressing paradoxes that arose in naive set theory. Consider Russell’s famous paradox: let $S = \{x \mid x \notin x\}$ be the set of all sets that do not contain themselves. This construction is paradoxical, leading to the contradiction $S \in S \iff S \notin S$. Type theory resolves such issues by working with **types** as primary objects rather than sets, and by restricting which constructions are well-formed.

Dependent type theory, the framework underlying Lean, extends basic type systems by allowing types to depend on values. For instance, one can define the type of vectors of length n , where n is itself a value. This capability makes dependent type theory particularly expressive for formalizing mathematics.

Various proof assistants have been developed based on different variants of type theory, including Agda, Coq, Idris, and Lean. Each system makes different design choices regarding which rules and features to include. Lean adopts the **Calculus of Inductive Constructions**, which extends the Calculus of Constructions, introduced in Coq, with **inductive types**. Inductive types allow for the definition of structures such as natural numbers, lists, and trees.

A fundamental design feature of Lean is its **universe hierarchy** of types, with **Prop** (the proposition type) as a distinguished universe at the bottom. The **Prop** universe exhibits two special properties: **impredicativity** and **proof irrelevance**. Proof irrelevance means that all proofs of the same proposition are considered equivalent. What matters is whether a proposition can be proven, not which specific proof is given. This separation between propositions (**Prop**) and data types (**Type**) was first introduced in N.G. de Bruijn’s **AUTOMATH system** (1967) ([Tho99]).

For our purposes, we do not need to delve deeply into the theoretical foundations; instead, we will introduce the relevant concepts as needed while working

with Lean. The practical aspects of writing proofs in Lean will be our primary focus, and the theoretical machinery will be explained only insofar as it aids understanding of how to formalize mathematics effectively.

1.1 Lean first steps

In the language of type theory, and by extension in Lean, we write $x : X$ to mean that x is a **term** of type X . For example, $2 : \mathbb{N}$ annotates 2 as a natural number, or more precisely, as a term of the natural number type. Lean has internally defined types such as `Nat` or `ℕ` (you can type `\Nat` to get the Unicode symbol). The command `#check` allows us to inspect the type of any expression, term or variable.

Example 1.2

```
#check 2 -- 2 : Nat
#check 2 + 2 -- 2 + 2 : Nat
```

*[Try this example in Lean Web Editor] By following the link, you can try out the code in your browser. Lean provides a dedicated **infoview** panel on the right side. Position your cursor after `#check 2`, and the infoview will display the output `2 : Nat`. This dynamic interaction, where the infoview responds to your cursor position, is what makes Lean an **interactive theorem prover**. As you move through your code, the infoview continuously updates, showing computations, type information, and proof states at each location.*

At first glance, one might be tempted to view the colon notation as analogous to the membership symbol \in from set theory, treating types as if they were sets. While this intuition can be helpful initially, type theory offers a fundamentally richer perspective. The crucial insight is the **Curry-Howard correspondence**, also known as the **propositions-as-types** principle. This correspondence establishes a deep connection between mathematical proofs and programs: **propositions correspond to types**, and **proofs correspond to terms** inhabiting those types. Under this interpretation, a term $x : X$ can be understood in two more ways:

- As a **computational object**: x is a program or data structure of type X
- As a **logical object**: x is a proof of the proposition X

Lean is a concrete realization of the propositions-as-types principle, proving a theorem, within the language, amounts to constructing a term of the appropriate type. When we write `theorem_name : Proposition`, we are declaring that `theorem_name` is a proof (term) of `Proposition` (type). For example, consider proving that $2 + 2 = 4$:

Example 1.3

```
theorem two_plus_two_eq_four : 2 + 2 = 4 := rfl
```

Lean’s syntax is designed to resemble the language of mathematics. Here, we use the `theorem` keyword to encapsulate our proof, of the statement/proposition $2 + 2 = 4$, giving it the name `two_plus_two_eq_four`. This allows us to reference and reuse this result later in our code. After the semicolon, `:`, we introduce the statement; $2 + 2 = 4$. The `:=` operator expects the proof term that establishes the theorem’s validity. The proof itself consists of a single term: `rfl` (short for **reflexivity**). This is a proof term that works by **definitional equality**, Lean’s kernel automatically reduces both sides of the equation to their normal (definitional) form and verifies they are identical. Since $2 + 2$ computes to 4, the proof succeeds immediately. We can now use this theorem in subsequent proofs. For instance:

```
example : 1 + 1 + 1 + 1 = 4 := two_plus_two_eq_four
```

Well, this example is simple enough for Lean to evaluate by itself: $1 + 1 + 1 + 1 = 2 + 2 = 4$ and conclude with `two_plus_two_eq_four`. Actually, `rfl` would solve the equation similarly, so this is just applying `rfl` again (it’s a bit of cheating). Here, I used `example`, which is handy for defining anonymous expressions for demonstration purposes. Before diving into the discussion, here is another keyword, `def`, used to introduce definitions and functions.

```
def addOne (n : Nat) : Nat := n + 1
```

This definition expects a natural number as its parameter, written `(n : Nat)` and returns a natural number. [Run in browser]

Let’s now turn to how logic is handled in Lean and how the Curry-Howard isomorphism is reflected concretely.

2 Logic and Proposition as Types

2.1 First Order Logic

Logic is the study of reasoning, branching into various systems. We refer to **classical logic** as the one that underpins much of traditional mathematics. It’s the logic of the ancient Greeks (not fair) and truth tables, and it remains used nowadays for pedagogical reasons. We first introduce **propositional logic**, which is the simplest form of classical logic. Later we will extend this to **predicate (or first-order) logic**, which includes **predicates** and **quantifiers**. In this setting, a **proposition** is a statement that is either true or false, and a **proof** is a logical argument that establishes the truth of a proposition. Propositions can be combined with logical **connectives** such as “and” (\wedge), “or” (\vee), “not” (\neg), “false” (\perp), “true” (\top) “implies” (\Rightarrow), and “if and only if” (\Leftrightarrow). These connectives allow the creation of complex or compound propositions.

Here how connectives are defined in Lean:

Example 2.1 (Logical connectives in Lean)

```
#check And (a b : Prop) : Prop
#check Or (a b : Prop) : Prop
#check True : Prop
#check False : Prop
#check Not (a : Prop) : Prop
#check Iff (a b : Prop) : Prop
```

Prop stands for proposition, and it is an essential component of Lean's type system. For now, we can think of it as a special type whose inhabitants are proofs; somewhat paradoxically, a type of types.

Logic is often formalized through a framework known as the **natural deduction system**, developed by Gentzen in the 1930s ([Wad15]). This approach brings logic closer to a computable, algorithmic system. It specifies rules for deriving **conclusions** from **premises** (assumptions from other propositions), called **inference rules**.

Example 2.2 (Deductive style rule) *Here is an hypothetical example of inference rule.*

$$\frac{P_1 \quad P_2 \quad \dots \quad P_n}{C}$$

Where the P_1, P_2, \dots, P_n , above the line, are hypothetical premises and, the hypothetical conclusion C is below the line.

The inference rules needed are:

- **Introduction rules** specify how to form compound propositions from simpler ones, and
- **Elimination rules** specify how to use compound propositions to derive information about their components.

Let's look at how we can define some connectives first using natural deduction.

Conjunction (\wedge)

Introduction Rule

$$\frac{A \quad B}{A \wedge B} \wedge\text{-Intro}$$

Elimination Rule

$$\frac{A \wedge B}{A} \wedge\text{-Elim}_1 \qquad \frac{A \wedge B}{B} \wedge\text{-Elim}_2$$

Disjunction (\vee)

Introduction Rule

$$\frac{A}{A \vee B} \vee\text{-Intro}_1 \qquad \frac{B}{A \vee B} \vee\text{-Intro}_2$$

Elimination (Proof by cases)

$$\frac{A \vee B \quad [A] \vdash C \quad [B] \vdash C}{C} \vee\text{-Elim}$$

Implication (\rightarrow)

Introduction Rule

$$\frac{[A] \vdash B}{A \rightarrow B} \rightarrow\text{-Intro}$$

Elimination (Modus Ponens)

$$\frac{A \rightarrow B \quad A}{B} \rightarrow\text{-Elim}$$

Notation 2.3 We use $A \vdash B$ (called *turnstile*) to designate a deduction of B from A . It is employed in Gentzen's **sequent calculus** ([GTL89]) and mostly used in type theory. The square brackets around a premise $[A]$ mean that the premise A is meant to be **discharged** at the conclusion. The classical example is the introduction rule for the implication connective. To prove an implication $A \rightarrow B$, we assume A (shown as $[A]$), derive B under this assumption, and then discharge the assumption A to conclude that $A \rightarrow B$ holds without the assumption. The turnstile is predominantly used in judgments and type theory with the meaning of “entails that”.

2.2 Primitive Types

Type theory employs this procedure too, by referring to deduction rules as **judgments**. A type judgment has the form $\Gamma \vdash t : T$, meaning: under **context** Γ (a list of typed variables), the term t has type T . Using formal inference rules in the type judgment system, such as **introduction** and **elimination** rules, we can construct new compound types from existing ones.

Example 2.4 (Judgment style rule)

$$\frac{\Gamma \vdash \quad p_1 : P_1 \quad p_2 : P_2 \quad \cdots \quad P_n}{C}$$

Technically, there are two more inference rules that we will not consider in this setting: **formation rules**, used to declare that a type is well-formed, and **computation rules**, which specify how a term will be evaluated. Moreover, without going too deep into the jargon, one specific judgment is $\Gamma \vdash A \equiv B$ type, which means “types A and B are **judgmentally (or definitionally) equal** in context Γ .” Similarly for terms, $\Gamma \vdash t_1 \equiv t_2 : A$ means “terms t_1 and t_2 are judgmentally equal of type A in context Γ .” In Lean, the operator `:=` stands for definitional equality and is used by the kernel to verify proof equality.

Let’s now construct new types from given types A and B .

Product Type As a fundamental example, $A \times B$ denotes the type of pairs (a, b) where $a : A$ and $b : B$, called the **product type**.

Introduction Rule (pairing)

$$\frac{a : A \quad b : B}{(a, b) : A \times B}$$

In Lean:

```
Prod.mk a b : Prod A B    -- or A × B
(a, b) : A × B
⟨a, b⟩ : A × B
```

Elimination Rules (projections)

$$\frac{p : A \times B}{\text{fst}(p) : A} \qquad \frac{p : A \times B}{\text{snd}(p) : B}$$

In Lean:

```
p.1 : A      -- or Prod.fst p
p.2 : B      -- or Prod.snd p
```

Sum Type The **sum type** $A + B$ (also called a coproduct or disjoint union) consists of values that are either of type A (tagged with `inl`) or of type B (tagged with `inr`).

Introduction Rules (injections)

$$\frac{a : A}{\text{inl}(a) : A + B} \qquad \frac{b : B}{\text{inr}(b) : A + B}$$

In Lean:

```
Sum.inl a : Sum A B    -- or A ⊕ B
Sum.inr b : Sum A B
```

Elimination Rule (case analysis)

$$\frac{p : A + B \quad f : (A \Rightarrow C) \quad g : (B \Rightarrow C)}{\text{cases}(p, f, g) : C}$$

In Lean, we can use the `cases`:

```
example (p : Sum A B) (f : A → C) (g : B → C) : C := by
  cases p with
  | inl x => f x
  | inr y => g y
```

Function Types The type of the form $A \rightarrow B$, used in the sum elimination rule represents functions from A to B .

Introduction Rule (lambda abstraction)

$$\frac{x : A \vdash \Phi : B}{\lambda x. \Phi : A \rightarrow B}$$

In Lean, lambda abstraction is written using `fun` or `λ`:

```
fun (x : A) => Φ : A → B
-- or using λ notation
λ (x : A) => Φ : A → B
-- Example: identity function
def id : A → A := fun x => x
-- or
def id : A → A := λ x => x
```

Elimination Rule (application)

$$\frac{f : A \rightarrow B \quad a : A}{f(a) : B}$$

In Lean, function application is written using juxtaposition:

```
example (f : A → B) (a : A) : B := f a
```

Functions are a primitive concept in type theory, and we provide a brief introduction here. We can **apply** a function $f : A \rightarrow B$ to an element $a : A$ to obtain an element of B , denoted $f(a)$. In type theory, it is common to omit the parentheses and write the application simply as $f a$.

There are two equivalent ways to construct function types: either by direct definition or by using λ -abstraction. Introducing a function by definition means that we introduce a function by giving it a name (let's say, f) and saying we define $f : A \rightarrow B$ by giving an equation

$$f(x) := \Phi \tag{1}$$

where x is a variable and Φ is an expression which may use x . In order for this to be valid, we have to check that $\Phi : B$ assuming $x : A$. Now we can compute $f(a)$ by replacing the variable x in Φ with a . As an example, consider the function $f : \mathbb{N} \rightarrow \mathbb{N}$ which is defined by $f(x) := x + x$. Then $f(2)$ is **definitionally equal** to $2 + 2$. If we don't want to introduce a name for the function, we can use **λ -abstraction**. Given an expression Φ of type B which may use $x : A$, as above, we write $\lambda(x : A).\Phi$ to indicate the same function defined by (1). Thus, we have

$$(\lambda(x : A).\Phi) : A \rightarrow B.$$

By convention, the “scope” of the variable binding “ λx .” is the entire rest of the expression, unless delimited with parentheses. Thus, for instance, $\lambda x.x + x$ should be parsed as $\lambda x.(x + x)$, not as $(\lambda x.x) + x$. Now a λ -abstraction is a function, so we can apply it to an argument $a : A$. We then have the following computation rule (β -reduction), which is a **definitional equality**:

$$(\lambda x.\Phi)(a) \equiv \Phi'$$

where Φ' is the expression Φ in which all occurrences of x have been replaced by a . Continuing the above example, we have $(\lambda x.x + x)(2) \equiv 2 + 2$. When performing calculations involving variables, we must carefully preserve the **binding structure** of expressions during substitution. Consider the function $f : \mathbb{N} \rightarrow (\mathbb{N} \rightarrow \mathbb{N})$ defined as:

$$f(x) := \lambda y.x + y$$

Suppose we have assumed $y : \mathbb{N}$ somewhere in our context. What is $f(y)$? A naive approach would replace x with y directly in the expression $\lambda y.x + y$, yielding $\lambda y.y + y$. However, this substitution is **semantically incorrect** because it causes **variable capture**: the free variable y (referring to our assumption) becomes bound by the λ -abstraction, fundamentally altering the expression's meaning. The correct approach uses **α -conversion** (variable renaming). Since bound variables have only local scope, we can consistently rename them while preserving binding structure. The expression $\lambda y.x + y$ is judgmentally equal to $\lambda z.x + z$ for any fresh variable z . Therefore:

$$f(y) \equiv \lambda z.y + z$$

This phenomenon parallels familiar mathematical practice: if $f(x) := \int_1^2 \frac{dt}{x-t}$, then $f(t)$ equals $\int_1^2 \frac{ds}{t-s}$, not the ill-defined $\int_1^2 \frac{dt}{t-t}$. Lambda abstractions bind dummy variables exactly as integrals do. For functions of multiple variables, we employ **currying** (named after mathematician Haskell Curry). Instead of using product types, we represent a two-argument function as a function returning another function. A function taking inputs $a : A$ and $b : B$ to produce output in C has type:

$$f : A \rightarrow (B \rightarrow C) \equiv A \rightarrow B \rightarrow C$$

where the arrow associates to the right by convention. Given $a : A$ and $b : B$, we apply f sequentially: first to a , then the result to b , obtaining $f(a)(b) : C$. To

simplify notation and avoid excessive parentheses, we adopt several conventions. We write $f(a)(b)$ as $f(a, b)$ for abbreviated application. Without parentheses entirely, $f a b$ means $(f a) b$ following left-associative application. For multi-parameter definitions, we write $f(x, y) := \Phi$ where $\Phi : C$ under assumptions $x : A$ and $y : B$. Using λ -abstraction, such definitions correspond to:

$$f := \lambda x. \lambda y. \Phi$$

Alternative notation using map symbols:

$$f := x \mapsto y \mapsto \Phi$$

This currying approach extends naturally to functions of three or more arguments, allowing us to represent any multi-argument function as a sequence of single-argument functions.

Example 2.5

```
def add : Nat -> (Nat -> Nat) := fun x => (fun y => x + y)
#eval add 3 4      -- Output: 7
```

Theoretically, lambda evaluation proceeds in steps:

$$\begin{aligned} \text{add } 3 \ 4 &\equiv (\text{add } 3) \ 4 \\ &\equiv ((\lambda x. \lambda y. x + y) \ 3) \ 4 \\ &\equiv ((\lambda y. x + y)[x := 3]) \ 4 \\ &\equiv (\lambda y. 3 + y) \ 4 \\ &\equiv (3 + y)[y := 4] \\ &\equiv 3 + 4 \\ &\equiv 7 \end{aligned}$$

2.3 Curry Howard isomorphism

We have been preparing for this argument, and the reader will have surely noticed a strong similarity when defining logical connectives using deduction rules; they are remarkably similar to types constructed using type judgments. For instance, function types can be seen as implications. This is not a coincidence, but rather a fundamental theorem first proven by Haskell Curry and William Howard. It forms the core of modern type theory and establishes a deep connection between logic, computation, and mathematics. The isomorphism states:

Propositions \leftrightarrow Types

Proofs \leftrightarrow Programs

Proof Normalization \leftrightarrow Program Evaluation

Implication ($P \Rightarrow Q$) corresponds to the **function type** ($P \rightarrow Q$). A proof of an implication is a function that transforms any proof of the premise into

a proof of the conclusion. **Conjunction** ($P \wedge Q$) corresponds to the **product type** ($P \times Q$). A proof of a conjunction consists of a pair containing proofs of both conjuncts. **Disjunction** ($P \vee Q$) corresponds to the **sum type** ($P + Q$). A proof of a disjunction is either a proof of the first disjunct or a proof of the second disjunct. Lean uses inference rules and type judgments as well as computing connectives using each related type. For instance, $A \wedge B$ can be represented as `And(A, B)` or $A \wedge B$. Its introduction rule is constructed by `And.intro _ _` or simply `<_, _>` (underscores are placeholders). The pair $A \wedge B$ can then be consumed using elimination rules `And.left` and `And.right`.

Example 2.6 *Let's look at a simple Lean example:*

```
example (ha : a) (hb : b) : (a ∧ b) := And.intro ha hb
```

This means: given a proof of a (ha) and a proof of b (hb), we can form a proof of (a ∧ b). And.intro is implemented as:

```
And.intro : p -> q -> (p ∧ q)
```

It says: if you give me a proof of p and a proof of q, then I return a proof of p ∧ q. We therefore conclude the proof by directly giving And.intro ha hb. Here is another way of writing the same statement:

```
example (ha : a) (hb : b) : And(a, b) := <ha, hb>
```

For a more concrete example, let's look at how proof normalization using a system of inference rules corresponds to computation in Lean. To reduce complexity of a **proof tree** in natural deduction, one follows a **top-down** approach, unfolding each component to be proved step by step.

Example 2.7 (Associativity of Conjunction) *We prove that $(A \wedge B) \wedge C$ implies $A \wedge (B \wedge C)$. First, from the assumption $(A \wedge B) \wedge C$, we can derive A:*

$$\frac{(A \wedge B) \wedge C}{\frac{A \wedge B}{A} \wedge E_1} \wedge E_1$$

Second, we can derive $B \wedge C$:

$$\frac{\frac{(A \wedge B) \wedge C}{\frac{A \wedge B}{B} \wedge E_2} \wedge E_1 \quad \frac{(A \wedge B) \wedge C}{C} \wedge E_2}{B \wedge C} \wedge I$$

Finally, combining these derivations we obtain $A \wedge (B \wedge C)$:

$$\frac{(A \wedge B) \wedge C \vdash A \quad (A \wedge B) \wedge C \vdash B \wedge C}{A \wedge (B \wedge C)} \wedge I$$

Example 2.8 (Lean Implementation) *Let us now implement the same proof in Lean.*

```

theorem and_associative (a b c : Prop) : (a ∧ b) ∧ c → a ∧ (b ∧ c) :=
  fun h : (a ∧ b) ∧ c →
    -- First, from the assumption (a ∧ b) ∧ c, we can derive a:
    have hab : a ∧ b := h.left
    have ha : a := hab.left
    -- Second, we can derive b ∧ c (here we only extract b and c and
    -- combine them in the next step)
    have hc : c := h.right
    have hb : b := hab.right
    -- Finally, combining these derivations we obtain a ∧ (b ∧ c)
    show a ∧ (b ∧ c) from ⟨ha, ⟨hb, hc⟩⟩

```

We introduce the `theorem` with the name `and_associative`. The type signature $(a \wedge b) \wedge c \rightarrow a \wedge (b \wedge c)$ represents our logical implication. Here, we construct the implication proof using a function (following the Curry Howard isomorphism) with the `fun` keyword. The `have` keyword introduces local lemmas within our proof scope, allowing us to break down complex reasoning into manageable intermediate steps, mirroring our natural deduction proof from before. Just before the keyword `show`, the info view displays the following context and goal:

```

a b c : Prop
h : (a ∧ b) ∧ c
hab : a ∧ b
ha : a
hc : c
hb : b
⊢ a ∧ b ∧ c

```

Finally, the `show` keyword explicitly states what we are proving and verifies that our provided term has the correct type. In this case, `show a ∧ (b ∧ c) from ⟨ha, ⟨hb, hc⟩⟩` asserts that we are constructing a proof of $a \wedge (b \wedge c)$ using the term $\langle ha, \langle hb, hc \rangle \rangle$. The `show` keyword serves two purposes: it makes the proof more readable by explicitly documenting what is being proved at this step, and it performs a type check to ensure the provided proof term matches the stated goal up to **definitional equality**. Two types are definitionally equal in Lean when they are identical after computation and unfolding of definitions; in other words, when Lean's type checker can mechanically verify they are the same without requiring additional proof steps. Here, the goal $\vdash a \wedge b \wedge c$ is definitionally equal to $a \wedge (b \wedge c)$ due to how conjunction associates, so `show` accepts this statement. If we had tried to use `show` with a type that was only **propositionally** equal (requiring a proof to establish equality) but not definitionally equal, Lean would reject it.

2.4 Predicate logic and dependency

To capture more complex mathematical ideas, we extend our system from propositional logic to **predicate logic**. A **predicate** is a statement or proposition that depends on a variable. In propositional logic we represent a proposition

simply by P . In predicate logic, this is generalized: a predicate is written as $P(a)$, where a is a variable. Notice that a predicate is just a function. This extension allows us to introduce **quantifiers**: \forall (“for all”) and \exists (“there exists”). These quantifiers express that a given formula holds either for every object or for at least one object, respectively. In Lean if α is any type, we can represent a predicate P on α as an object of type $\alpha \rightarrow \mathbf{Prop}$. Thus given an $x : \alpha$ (an element with type α) $P(x) : \mathbf{Prop}$ would be representative of a proposition holding for x .

We can give an informal reading of the quantifiers as infinite logical operations:

$$\begin{aligned}\forall x. P(x) &\equiv P(a) \wedge P(b) \wedge P(c) \wedge \dots \\ \exists x. P(x) &\equiv P(a) \vee P(b) \vee P(c) \vee \dots\end{aligned}$$

The expression $\forall x. P(x)$ can be understood as a generalized form of conjunction. It expresses that P holds for all possible values of x . Similarly, $\exists x. P(x)$ is a generalized disjunction, expressing that P holds for at least one value of x . Under the Curry-Howard isomorphism, universal quantifiers correspond to **dependent function types** (also called Pi types, written Π), while existential quantifiers correspond to **dependent pair types** (also called Sigma types, written Σ). These are constructs from dependent type theory, which provides a way to interpret predicates or, more generally, types depending on some data or variable. This time we are not going to involve deduction rules or type judgments. Instead, we will extend the isomorphism to quantifiers directly by presenting the Lean syntax.

Example 2.9 (Quantifiers in Lean) *Lean expresses quantifiers as follows:*

```

∀ (x :  $\alpha$ ), P x
forall (x :  $\alpha$ ), P x
-- Equivalently, using Pi types
Π (x :  $\alpha$ ), P x

∃ x :  $\alpha$ , p
exists (λ x :  $\alpha$  => p)
-- Equivalently, using Sigma types
Σ x :  $\alpha$ , p

```

Example 2.10 (Universal introduction in Lean) *The **universal introduction rule** allows us to prove $\forall x, P(x)$ by proving $P(x)$ for an **arbitrary** x . In Lean, this corresponds to lambda abstraction (constructing a function):*

```
example :  $\forall n : \text{Nat}, n \geq 0 :=$ 
  fun n => Nat.zero_le n
```

Example 2.11 (Universal elimination in Lean) *The **universal elimination rule** allows us to instantiate a universally quantified statement with a specific value. In Lean, this is simply function application:*

```
example (h :  $\forall n : \text{Nat}, n \geq 0$ ) :  $5 \geq 0 :=$ 
  h 5
```

Example 2.12 (Existential introduction in Lean) *When introducing an **existential** proof, we need a **pair** consisting of a witness and a proof that this witness satisfies the statement.*

```
example (x : Nat) (h :  $x > 0$ ) :  $\exists y, y < x :=$ 
  ⟨0, h⟩
```

Notice that $\langle 0, h \rangle$ is a product type holding data (the witness 0) and a proof that it satisfies the property.

Example 2.13 (Existential elimination in Lean) *The **existential elimination rule** (`Exists.elim`) allows us to prove a proposition Q from $\exists x, P(x)$ by showing that Q follows from $P(w)$ for an **arbitrary** value w . The existential quantifier can be interpreted as an infinite disjunction, so existential elimination naturally corresponds to a **proof by cases** (with a single case). In Lean, this is done using **pattern matching** with `cases`:*

```

example (h :  $\exists n : \text{Nat}, n > 0$ ) :  $\exists n : \text{Nat}, n > 0$  := by
  cases h with
  / intro witness proof => (witness, proof)

```

2.5 Constructive Mathematics

Mathematicians have traditionally worked within **classical logic**, using **sets** as the primary means of structuring mathematical objects. In contrast, **type theory** does not take sets as its primitive notion, nor is it built by first applying logic and then adding structure. Instead, logic is internal to type theory and is based on **constructive** (or **intuitionistic**) logic, introduced by Brouwer and formalized by Heyting (see, e.g., [GTL89]).

A major point of departure from classical logic is that, in constructive logic, statements cannot simply be classified as true or false; their truth depends on whether a proof exists. There are many conjectures, such as the Riemann Hypothesis, for which we do not yet know whether a proof or disproof exists, so we cannot say whether they are true or false. Consequently, constructive logic does not universally accept principles such as the **axiom of choice** or the **law of excluded middle** (every proposition is either true or false) as axioms. As a consequence, proof by contradiction does not work in this setting without additional justification.

Constructive logic emphasizes that a statement is only considered true if we can explicitly construct a proof or provide a **witness** for it. This is what makes constructive mathematics inherently **computable**.

We already touched on this concept in the previous section. In particular, we presented the logical connectives via the Brouwer–Heyting–Kolmogorov (BHK) interpretation. Following this interpretation, negation is not a primitive type but is instead constructed as a function $\text{Prop} \rightarrow \text{False}$. We also emphasized that, constructively, a proof of existence consists of a pair: a witness together with a proof that the stated property holds for that witness.

Example 2.14 (Constructive existence proof) *We give a constructive proof in Lean that there exist natural numbers a and b such that $a + b = 7$:*

```

example :  $\exists a b : \text{Nat}, a + b = 7$  := by
  use 3, 4

```

The `use` tactic (from `Mathlib`) provides explicit witnesses: $a = 3$ and $b = 4$. Lean then automatically evaluates the expression and verifies that $3 + 4 = 7$. This example is simple enough for Lean to complete the proof automatically.

In classical mathematics, one might attempt a proof by contradiction. However, this approach is not directly accepted in constructive mathematics, as it doesn't provide explicit witnesses for the claimed objects. Nonetheless, while constructive at its core, Lean allows users to invoke classical principles, such

as contraposition or proof by contradiction, through tactics like `exfalso` or by importing `Classical`.

Example 2.15 (Reasoning from false) *Here is an example of deriving any proposition from a contradiction:*

```
example (p : Prop) (h : False) : p := by
  exfalso
  exact h
```

*This example takes a proposition p to prove and a false hypothesis h . The `exfalso` tactic transforms the goal into $\vdash \text{False}$, meaning we now need to derive a contradiction. Since we already have a false hypothesis h , we can provide it using the `exact` tactic. This principle is known as *ex falso quodlibet* (from falsehood, anything follows).*

3 Describing and use properties

It is interesting to note that a relation can be expressed as a function: $R : \alpha \rightarrow \alpha \rightarrow \text{Prop}$. Similarly, when defining a predicate ($P : \alpha \rightarrow \text{Prop}$) we must first declare $\alpha : \text{Type}$ to be some arbitrary type. This is what is called **polymorphism**, more specifically **parametrical polymorphism**. A canonical example is the identity function, written as $\alpha \rightarrow \alpha$, where α is a type variable. It has the same type for both its domain and codomain, this means it can be applied to booleans (returning a boolean), numbers (returning a number), functions (returning a function), and so on. In the same spirit, we can define a transitivity property of a relation as follows:

```
def Transitive (α : Type) (R : α → α → Prop) : Prop :=
  ∀ x y z, R x y → R y z → R x z
```

To use `Transitive`, we must provide both the type α and the relation itself. For example, here is a proof of transitivity for the less-than relation on \mathbb{N} (in Lean `Nat` or \mathbb{N}):

```
theorem le_trans_proof : Transitive Nat (· ≤ · : Nat → Nat → Prop) :=
  fun x y z h1 h2 => Nat.le_trans h1 h2 -- this lemma is provided by Lean
```

Looking at this code, we immediately notice that explicitly passing the type argument `Nat` is somewhat repetitive. Lean allows us to omit it by letting the type inference mechanism fill it in automatically. This is achieved by using **implicit arguments** with curly brackets:

```
def Transitive {α : Type} (R : α → α → Prop) : Prop :=
  ∀ x y z, R x y → R y z → R x z
theorem le_trans_proof : Transitive (· ≤ · : Nat → Nat → Prop) :=
  fun x y z h1 h2 => Nat.le_trans h1 h2
```

Lean's type inference system is quite powerful: in many cases, types can be completely inferred without explicit annotations.

Example 3.1 (Type Inference in Lean)

```
def double (n : Nat) := n + n
-- Lean infers return type is Nat because n : Nat and + : Nat → Nat →
  Nat
def id {α : Type} (x : α) : α := x
#check id 5          -- Lean infers α = Nat
#check id "hello"    -- Lean infers α = String
```

Let us now revisit the transitivity proof, but this time for the less-than-equal relation on the rational numbers (\mathbf{Rat} or \mathbb{Q}) instead.

```
import Mathlib

theorem rat_le_trans : Transitive (· ≤ · : Rat → Rat → Prop) :=
  fun _ _ h1 h2 => Rat.le_trans h1 h2
```

Here, `Rat.le_trans` is the transitivity lemma for \leq on rational numbers, provided by `Mathlib`. We import `Mathlib` to access `Rat` and `le_trans`. `Mathlib` is the community-driven mathematical library for Lean, containing a large body of formalized mathematics and ongoing development. It is the defacto standard library for both programming and proving in Lean [Com20], we will dig into it as we go along. Notice that we used a function to discharge the universal quantifiers required by transitivity. The underscores indicate unnamed variables that we do not use later. If we had named them, say `x y z`, then: `h1` would be a proof of $x \leq y$, `h2` would be a proof of $y \leq z$, and `Rat.le_trans h1 h2` produces a proof of $x \leq z$. The `Transitive` definition is imported from `Mathlib` and similarly defined as before.

Example 3.2 *The code can be made more readable using **tactic mode**. In this mode, you use tactics, commands provided by Lean or defined by users, to carry out proof steps succinctly, avoid code repetition, and automate common patterns. This often yields shorter, clearer proofs than writing the full term by hand.*

```
import Mathlib

theorem rat_le_trans : Transitive (· ≤ · : Rat → Rat → Prop) := by
  intro x y z hxy hyz
  exact Rat.le_trans hxy hyz
```

This proof performs the same steps but is much easier to read. Using `by` we enter Lean's tactic mode. Move your cursor just before `by`. The goal is initially displayed as $\vdash \text{Transitive } \text{fun } x1\ x2 \mapsto x1 \leq x2$. The tactic `intro` is mainly used to introduce variables and hypotheses corresponding to universal quantifiers and assumptions into the context (essentially deconstructing universal quantifiers and implications). Now position your cursor just before `exact` and observe the info view again. The goal is now $\vdash x \leq z$, with the context showing the variables and hypotheses introduced by the previous tactic. The `exact` tactic closes the goal by supplying the term `Rat.le_trans hxy hyz` that exactly matches the goal (the specification of `Transitive`). You can hover over each tactic to see its definition and documentation.

3.1 Exploring Mathlib (The Rat structure)

In these examples we cheated and have used predefined lemmas such as `Nat.le_trans` and `Rat.le_trans`, just to simplify the presentation. We can now dig into the implementation of these lemmas. Let's look at the source code of `Rat.le_trans`. The `Mathlib 4` documentation website is at https://leanprover-community.github.io/mathlib4_docs, and the documentation for `Rat.le_trans` is at https://leanprover-community.github.io/mathlib4_docs/Rat.html#Rat.le_trans.

[//leanprover-community.github.io/mathlib4_docs/Mathlib/Algebra/Order/Ring/Unbundled/Rat.html#Rat.le_trans](https://leanprover-community.github.io/mathlib4_docs/Mathlib/Algebra/Order/Ring/Unbundled/Rat.html#Rat.le_trans). Click the "source" link there to jump to the implementation in the Mathlib repository. In editors like VS Code you can also jump directly to the definition (Ctrl+click; Cmd+click on macOS). Another way to check source code is by using `#print Rat.le_trans`.

```
variable (a b c : Rat)
protected lemma le_trans (hab : a ≤ b) (hbc : b ≤ c) : a ≤ c := by
  rw [Rat.le_iff_sub_nonneg] at hab hbc
  have := Rat.add_nonneg hab hbc
  simp_rw [sub_eq_add_neg, add_left_comm (b + -a) c (-b), add_comm (b +
    -a) (-b), add_left_comm (-b) b (-a), add_comm (-b) (-a),
    add_neg_cancel_comm_assoc, ← sub_eq_add_neg] at this
  rwa [Rat.le_iff_sub_nonneg]
```

The proof uses several tactics and lemmas from Mathlib. The `rw` or `rewrite` tactic is very common and syntactically similar to the mathematical practice of rewriting an expression using an equality. In this case, with `at`, we use it to rewrite the hypotheses `hab` and `hbc` using another Mathlib's lemma `Rat.le_iff_sub_nonneg`, which states that for any two rational numbers x and y , $x \leq y$ is equivalent to $0 \leq y - x$. Thus we now have the hypotheses transformed to :

```
hab : 0 ≤ b - a
hbc : 0 ≤ c - b
```

The `have` tactic introduces an intermediate result. If you omit a name, Lean assigns it the default name `this`. In our situation, from `hab : a ≤ b` and `hbc : b ≤ c` we can derive that `b - a` and `c - b` are nonnegative, hence their sum is nonnegative:

```
this : 0 ≤ b - a + (c - b)
```

The most involved step uses `simp_rw` to simplify the expression via a sequence of other existing Mathlib's lemmas. The tactic `simp_rw` is a variant of `simp`: it performs rewriting using the simp set (and any lemmas you provide), applying the rules in order and in the given direction. Lemmas that `simp` can use are typically marked with the `@[simp]` attribute. This is particularly useful for simplifying algebraic expressions and equations. After these simplifications we obtain:

```
this : 0 ≤ c - a
```

Clearly, the proof relies mostly on `Rat.add_nonneg`. Its source code is fairly involved and uses advanced features that are beyond our current scope. Nevertheless, it highlights an important aspect of formal mathematics in Mathlib. Mathlib defines `Rat` as an instance of a linear ordered field, implemented via a normalized fraction representation: a pair of integers (numerator and denominator) with positive denominator and coprime numerator and denominator [Lea25b]. To achieve this, it uses a **structure**. In Lean, a structure is a dependent record (or product type) type used to group together related fields or properties as a single data type. Unlike ordinary records, the type of later fields

may depend on the values of earlier ones. Defining a structure automatically introduces a constructor (usually `mk`) and projection functions that retrieve (deconstruct) the values of its fields. Structures may also include proofs expressing properties that the fields must satisfy.

```

structure Rat where
  /-- Constructs a rational number from components.
  We rename the constructor to 'mk' to avoid a clash with the smart
  constructor. -/
  mk' ::
  /-- The numerator of the rational number is an integer. -/
  num : Int
  /-- The denominator of the rational number is a natural number. -/
  den : Nat := 1
  /-- The denominator is nonzero. -/
  den_nz : den ≠ 0 := by decide
  /-- The numerator and denominator are coprime: it is in "reduced
  form". -/
  reduced : num.natAbs.Coprime den := by decide

```

In order to work with rational numbers in Mathlib, we use the `Rat.mk'` constructor to create a rational number from its numerator and denominator, if omitted the default would be `Rat.mk`. The fields `den_nz` and `reduced` are proofs that the denominator is nonzero and that the numerator and denominator are coprime, respectively. These proofs are automatically generated by Lean's `decide` tactic, which can solve certain decidable propositions (to be discussed in the next section).

Example 3.3 *Here is how we can define and manipulate rational numbers in Lean.*

```

def half : Rat := Rat.mk' 1 2
def third : Rat := Rat.mk' 1 3

```

When working with rational numbers, or more generally with structures, we must provide the required proofs as arguments to the constructor (or Lean must be able to ensure them). For instance `Rat.mk' 1 0` or `Rat.mk' 2 6` would be rejected. In the case of rationals, Mathlib unfolds the definition through `Rat.numDenCasesOn`. This principle states that, to prove a property of an arbitrary rational number, it suffices to consider numbers of the form n / d in canonical (normalized) form, with $d > 0$ and $\gcd n d = 1$. This reduction allows mathlib to transform proofs about \mathbb{Q} into proofs about \mathbb{Z} and \mathbb{N} , and then lift the result back to rationals.

Example 3.4 *We present a simplified implementation of addition non-negativity for rationals (`Rat.add_nonneg`), maintaining a similar approach: projecting everything to the natural numbers and integers first. To illustrate the proof technique clearly, we avoid using existing lemmas from the `Rat` module in Mathlib. Mathlib is indeed organized into modules by mathematical domain (e.g., `Nat`, `Int`, `Rat`).*

We start by defining helper lemmas needed in the main proof. Given a natural number (which in this case represents the denominator of a rational number) that is not equal to zero, we prove it must be positive. This follows directly by applying the Mathlib lemma `Nat.pos_of_ne_zero`:

```
import Mathlib
```

```
lemma nat_ne_zero_pos (den : ℕ) (h_den_nz : den ≠ 0) : 0 < den :=
  Nat.pos_of_ne_zero h_den_nz
```

The naming convention follows Mathlib best practices aiming to be descriptive by indicating types and properties involved.

The following lemma is slightly more involved. It states that if a rational number (num / den) is non-negative, then its numerator must also be non-negative:

```
lemma rat_num_nonneg {num : ℤ} {den : ℕ} (hden_pos : 0 < den)
  (h : (0 : ℚ) ≤ num / den) : 0 ≤ num := by
  contrapose! h
  have hden_pos_to_rat : (0 : ℚ) < den := Nat.cast_pos.mpr hden_pos
  have hnum_neg_to_rat : num < (0 : ℚ) := Int.cast_lt.mpr h
  exact div_neg_of_neg_of_pos hnum_neg_to_rat hden_pos_to_rat
```

The lemma requires the denominator to be positive as well as the non-negativity of the rational number, expressed as num / den where the types of num and den are inferred. First, notice the type annotation $(0 : \mathbb{Q})$. This explicit type annotation on zero forces the entire equation to be casted into rational numbers. Without this annotation, Lean would infer 0 as a natural number by default. However, since the main theorem we are proving concerns rational numbers, we must ensure all comparisons occur in \mathbb{Q} . The tactic `contrapose!` does what you might expect: it proves a statement by contraposition. According to the documentation:

- `contrapose` turns a goal $P \rightarrow Q$ into $\neg Q \rightarrow \neg P$
- `contrapose!` turns a goal $P \rightarrow Q$ into $\neg Q \rightarrow \neg P$ and pushes negations inside P and Q using `push_neg`
- `contrapose h` first reverts the local assumption h , then uses `contrapose` and `intro h`
- `contrapose! h` first reverts the local assumption h , then uses `contrapose!` and `intro h`

In our case, `contrapose! h` transforms the goal from proving $0 \leq \text{num}$ to assuming $\text{num} < 0$ and proving $\text{num} / \text{den} < 0$. We then introduce two local hypotheses. The first, `hden_pos_to_rat`, proves that the denominator is positive when cast to rationals, using `Nat.cast_pos`. The suffix `.mpr` selects the “modus ponens reverse” direction of the biconditional (the \leftarrow direction of the \leftrightarrow). Next, we introduce `hnum_neg_to_rat`, which expresses that the numerator is negative when cast to rationals, using `Int.cast_lt` with `.mpr` again. Finally, we apply `div_neg_of_neg_of_pos`, which states that dividing a negative number by a positive number yields a negative result, thus completing the proof by contraposition. Note that we are allowing ourselves to use existing lemmas from Mathlib, such as `div_neg_of_neg_of_pos` from the `Field` module, but not from the `Rat` module, to keep the presentation clear and focused on the main proof techniques.

Now we can prove the main result:

```

lemma rat_add_nonneg (a b : Rat) : 0 ≤ a → 0 ≤ b → 0 ≤ a + b := by

  intro ha hb
  cases a with / div a_num a_den a_den_nz a_cop =>
  cases b with / div b_num b_den b_den_nz b_cop =>
  -- Goal: ⊢ 0 ≤ ↑a_num / ↑a_den + ↑b_num / ↑b_den
  rw[div_add_div] -- applies the addition formula requiring two new goals
  · sorry
  · sorry
  · sorry

```

We first introduce the two hypotheses `ha` and `hb` into the context using `intro`. As mentioned earlier, a structure can be viewed as a product type or a record type with a single constructor. The tactic `cases a with` exposes the fields of `Rat`: the numerator (`a_num`), denominator (`a_den`), the proof that the denominator is non-zero (`a_den_nz`), and the coprimality condition (`a_cop`). Notice how the goal transforms the rationals `a` and `b` into:

$$\vdash 0 \leq \uparrow a_num / \uparrow a_den + \uparrow b_num / \uparrow b_den$$

where \uparrow denotes type coercion from \mathbb{Z} or \mathbb{N} to \mathbb{Q} . Now we rewrite the goal using `rw [div_add_div]`, a theorem from the `Field` module, which applies the addition formula for division. Let us briefly examine the source code of this theorem:

```
variable [Semifield K] {a b d : K}
```

```
theorem div_add_div (a : K) (c : K) (hb : b ≠ 0) (hd : d ≠ 0) :
  a / b + c / d = (a * d + b * c) / (b * d) := ...
```

The type `K` here is assumed to be a `Semifield`. The `variable` keyword is a way to declare parameters that are potentially used across multiple theorems or definitions. We will explore Lean’s powerful algebraic hierarchy and the meaning of the square brackets `[]` in a later section. Using this rewrite is particularly time-saving, since otherwise one would have to establish the well-definedness of rational addition in terms of the underlying structure (a non-trivial task). This theorem requires proofs `(hb : b ≠ 0)` and `(hd : d ≠ 0)`, generating two additional side goals. We handle each goal separately using the focusing bullet `·`. The first bullet addresses the main goal (proving the sum is non-negative), while the subsequent bullets discharge the non-zero denominator conditions. I have omitted the actual proofs, here, using `sorry`, which we haven’t mentioned before. `sorry` is a useful feature of Lean that tells the system to accept an incomplete proof for the time being, allowing you to continue development without proving every detail immediately. We can now tackle the remaining goals:


```

· -- Goal:  $\vdash 0 \leq (\uparrow a\_num * \uparrow b\_den + \uparrow a\_den * \uparrow b\_num) / (\uparrow a\_den * \uparrow b\_den)$ 
have hnum_nonneg :  $(0 : \mathbb{Q}) \leq a\_num * b\_den + a\_den * b\_num := by$ 
  have ha_num_nonneg := by
    have ha_den_pos := nat_ne_zero_pos a_den a_den_nz
    exact rat_num_nonneg ha_den_pos ha
  have hb_num_nonneg := by
    have hb_den_pos := nat_ne_zero_pos b_den b_den_nz
    exact rat_num_nonneg hb_den_pos hb
  apply add_nonneg -- works for any OrderedAddCommMonoid
· apply mul_nonneg -- works for any OrderedSemiring
  · exact Int.cast_nonneg.mpr ha_num_nonneg
  · exact Nat.cast_nonneg b_den
· apply mul_nonneg
  · exact Nat.cast_nonneg a_den
  · exact Int.cast_nonneg.mpr hb_num_nonneg

have hden_nonneg :  $(0 : \mathbb{Q}) \leq a\_den * b\_den := by$ 
  rw [← Nat.cast_mul]
  exact Nat.cast_nonneg (a_den * b_den)
exact div_nonneg hnum_nonneg hden_nonneg

· exact Nat.cast_ne_zero.mpr a_den_nz -- Goal:  $\vdash \uparrow a\_den \neq 0$ 
· exact Nat.cast_ne_zero.mpr b_den_nz -- Goal:  $\vdash \uparrow b\_den \neq 0$ 

```

We introduce two key hypotheses, `hnum_nonneg` and `hden_nonneg`, which will be required by `div_nonneg` from the `GroupWithZero` module. This lemma provides us with a term that directly validates our statement. Note that `div_nonneg` is a generalized lemma that applies not only to rational numbers but to all ordered groups with zero that are also partially ordered. The hypothesis `hnum_nonneg` proves that the numerator is non-negative by working with the coerced expressions in \mathbb{Q} . It uses `add_nonneg` and `mul_nonneg`, which are general theorems that work for any ordered additive commutative monoid and ordered semiring, respectively. The actual reasoning is done using integer-related theorems (via `Int.cast_nonneg`) for the numerators and natural number theorems (via `Nat.cast_nonneg`) for the denominators. The hypothesis `hden_nonneg` proves that the denominator is non-negative by working entirely with natural numbers. We use the rewrite `rw [← Nat.cast_mul]`, which moves the coercion (in this case from \mathbb{N} to \mathbb{Q}) inside the multiplication: $\uparrow(m * n) = \uparrow m * \uparrow n$. The `←` symbol means that we want the transformation from right to left (i.e., we apply the equality in reverse to move the cast inward). Type casts and coercions require these kinds of rewrite rules, not only for multiplication but also for addition and other operations, and similarly for \mathbb{Z} or other numerical types. These lemmas, such as `Nat.cast_mul`, `Int.cast_add`, etc., ensure that algebraic operations commute with type coercions.

3.2 Coercions and Type Casting

We extensively used type casting and coercions in this proof, which requires some explanation [LM20]. Lean’s type system lacks subtyping, means that types like \mathbb{N} , \mathbb{Z} , and \mathbb{Q} are distinct and do not have a subtype relationship. In order to translate between these types, we need to use explicit type casts or rely on automatic coercions. For example, natural numbers (\mathbb{N}) can be coerced to integers (\mathbb{Z}), and integers can be coerced to rational numbers (\mathbb{Q}). Casting and coercion are related but distinct concepts:

- **Casting** refers to the explicit conversion of a value from one type to another, typically using functions like `Int.cast` or `Nat.cast`. These functions have accompanying lemmas that preserve properties across type conversions, such as `Int.cast_lt` and `Nat.cast_pos`.
- **Coercion**, on the other hand, is a more general mechanism that allows Lean to automatically convert between types when needed. More generally, in expressions like $x + y$ where x and y are of different types, Lean will automatically coerce them to a common type. For example, if $x : \mathbb{N}$ and $y : \mathbb{Z}$, then x will be coerced to \mathbb{Z} .

The notation \uparrow denotes an explicit coercion (in between cast and coercion). To illustrate the expected behavior of coercion simplification, consider the expression $\uparrow m + \uparrow n < (10 : \mathbb{Z})$, where $m, n : \mathbb{N}$ are cast to \mathbb{Z} . The expected normal form is $m + n < (10 : \mathbb{N})$, since $+$, $<$, and the numeral 10 are polymorphic (i.e., they can work with any numerical type such as \mathbb{Z} or \mathbb{N}). The simplification should proceed as follows:

1. Replace the numeral on the right with the cast of a natural number: $\uparrow m + \uparrow n < \uparrow(10 : \mathbb{N})$
2. Factor \uparrow to the outside on the left: $\uparrow(m + n) < \uparrow(10 : \mathbb{N})$
3. Eliminate both casts to obtain an inequality over \mathbb{N} : $m + n < (10 : \mathbb{N})$

Lean provides tactics like `norm_cast` to simplify expressions involving such coercions. The `norm_cast` tactic normalizes casts by pushing them outward and eliminating redundant coercions, often simplifying proofs significantly by reducing goals to their “native” types.

3.3 Type Classes and Algebraic Hierarchy in Lean

In our proof of `rat_add_nonneg`, we used many generalized lemmas from Mathlib, such as `add_nonneg`, `mul_nonneg`, and `div_nonneg`, which apply to a wide range of types beyond just rational numbers. Similarly, in our earlier work with natural numbers, we used `Nat.le_trans`, a theorem specifically for natural numbers that is part of Lean’s core library (`lean/Init/Prelude.lean`). Mathlib is built on top of this base library. However, the transitivity property holds not only for naturals but also for integers, reals, and, in fact, for any partially ordered set.

Rather than duplicating this theorem for each type, Mathlib provides a general lemma `le_trans` that works for any type α endowed with a partial ordering. This is achieved through **type classes**, Lean’s mechanism for defining and working with abstract algebraic structures in an ad hoc polymorphic manner. Type classes provide a powerful and flexible way to specify properties and operations that can be shared across different types, thereby enabling polymorphism and code reuse. Ad hoc polymorphism arises when a function or operator is defined over several distinct types, with behavior that varies depending on the type. A standard example [WB89] is overloaded multiplication: the same symbol `*` denotes multiplication of integers (e.g., `3 * 3`) and of floating-point numbers (e.g., `3.14 * 3.14`). By contrast, parametric polymorphism occurs when a function is defined over a range of types but acts uniformly on each of them. For instance, the `List.length` function applies in the same way to a list of integers and to a list of floating-point numbers.

Under the hood, a type class is a structure. An important aspect of structures, and hence type classes, is that they support hierarchy and composition through inheritance. For example, mathematically, a monoid is a semigroup with an identity element, and a group is a monoid with inverses. In Lean, we can express this by defining a `Monoid` structure that extends the `Semigroup` structure, and a `Group` structure that extends the `Monoid` structure using the `extends` keyword:

```
-- A semigroup has an associative binary operation
structure Semigroup (α : Type*) where
  mul : α → α → α
  mul_assoc : ∀ a b c : α, mul (mul a b) c = mul a (mul b c)

-- A monoid extends semigroup with an identity element
structure Monoid (α : Type*) extends Semigroup α where
  one : α
  one_mul : ∀ a : α, mul one a = a
  mul_one : ∀ a : α, mul a one = a

-- A group extends monoid with inverses
structure Group (α : Type*) extends Monoid α where
  inv : α → α
  mul_left_inv : ∀ a : α, mul (inv a) a = one
```

The symbol `*` in `(α : Type*)` indicates a universe variable (we will discuss universes later). Sometimes, to avoid inconsistencies between types (such as Girard’s paradox), universes must be specified explicitly. This is an example of universe polymorphism. Thus we have now seen all the polymorphism flavors in Lean: parametric, ad hoc, and universe polymorphism.

Type classes are defined using the `class` keyword, which is syntactic sugar for defining a structure. Thus, the previous example can be rewritten using type classes:

```

-- A semigroup has an associative binary operation
class Semigroup ( $\alpha$  : Type*) where
  mul :  $\alpha \rightarrow \alpha \rightarrow \alpha$ 
  mul_assoc :  $\forall a b c : \alpha, \text{mul} (\text{mul } a b) c = \text{mul } a (\text{mul } b c)$ 

-- A monoid extends semigroup with an identity element
class Monoid ( $\alpha$  : Type*) extends Semigroup  $\alpha$  where
  one :  $\alpha$ 
  one_mul :  $\forall a : \alpha, \text{mul one } a = a$ 
  mul_one :  $\forall a : \alpha, \text{mul } a \text{ one} = a$ 

-- A group extends monoid with inverses
class Group ( $\alpha$  : Type*) extends Monoid  $\alpha$  where
  inv :  $\alpha \rightarrow \alpha$ 
  mul_left_inv :  $\forall a : \alpha, \text{mul} (\text{inv } a) a = \text{one}$ 

```

The main difference is that type classes support **instance resolution**. We use the keyword **instance** to declare that a particular type is an instance of a type class, which inherits the properties and operations defined in the type class. Instances can be automatically inferred by Lean's type inference system, allowing for concise and expressive code. For example, we can declare that \mathbb{Z} is a group under addition:

```

instance : Group  $\mathbb{Z}$  where
  mul := Int.add
  one := 0
  inv := Int.neg
  mul_assoc := Int.add_assoc
  one_mul := Int.zero_add
  mul_one := Int.add_zero
  mul_left_inv := Int.neg_add_cancel

```

Now, any theorem proven for an arbitrary Group α automatically applies to \mathbb{Z} without any additional work. This mechanism is particularly useful for defining and working with order structures like preorders and partial orders. Mathematically, a preorder consists of a set P and a binary relation \leq on P that is reflexive and transitive [Lea25a].

```

-- A preorder is a reflexive, transitive relation ' $\leq$ ' with ' $<$ ' defined
   in terms of ' $\leq$ '
class Preorder ( $\alpha$  : Type*) extends LE  $\alpha$ , LT  $\alpha$  where
  le_refl :  $\forall a : \alpha, a \leq a$ 
  le_trans :  $\forall a b c : \alpha, a \leq b \rightarrow b \leq c \rightarrow a \leq c$ 
  lt := fun a b =>  $a \leq b \wedge \neg b \leq a$ 
  lt_iff_le_not_ge :  $\forall a b : \alpha, a < b \leftrightarrow a \leq b \wedge \neg b \leq a$  := by intros;
    rfl

instance [Preorder  $\alpha$ ] : Lean.Grind.Preorder  $\alpha$  where
  le_refl := Preorder.le_refl
  le_trans := Preorder.le_trans _ _ _
  lt_iff_le_not_ge := Preorder.lt_iff_le_not_ge _ _

```

Listing 1: Preorder Type Class in Lean

The `class Preorder` declares a type class over a type α , bundling the \leq and $<$ relations (inherited via `extends LE α , LT α`) with the preorder axioms: reflexivity (`le_refl`) and transitivity (`le_trans`). The field `lt` provides a default definition of strict inequality in terms of \leq , and the theorem `lt_iff_le_not_ge` characterizes this relationship, proved automatically via reflexivity (`by intros; rfl`). The `instance` declaration connects the `Preorder` class to Lean's `Grind` tactic automation, which allows automatic reasoning with preorder properties during proof search. Returning to our rational number proof, this explains why lemmas like `add_nonneg` and `mul_nonneg` work seamlessly: \mathbb{Q} is an instance of `OrderedSemiring`, which extends `Preorder` and other algebraic structures, automatically providing all their theorems.

We have roughly seen how Lean constructively builds the rational numbers from naturals and integers. Using the power of structures and type classes, Mathlib generalizes these concepts further to develop rich mathematical theories. However, when dealing with real numbers, the approach taken includes the use of the axiom of choice, which, as we discussed in constructive mathematics, is not directly accepted constructively. When constructive methods are insufficient, Lean provides classical axioms through the `Classical` module. For instance, the law of excluded middle become available:

```

open Classical

example (p : Prop) : p  $\vee$   $\neg p$  := em p

```

Using classical axioms comes at a cost: definitions and theorems that depend on them must be marked `noncomputable`. For example, many operations on real numbers, such as computing the sine function requires this marker:

```

noncomputable def realSin (x :  $\mathbb{R}$ ) :  $\mathbb{R}$  := Real.sin x

```

In the next section, I will present an example of formalization that requires working with real numbers as well as topological spaces, which provide the foundational tools for real analysis concepts like continuity and convergence. Topological spaces in Mathlib are built upon the concept of open sets using the

`TopologicalSpace` type class, which can be extended to define metric spaces and normed spaces. This hierarchical organization allows definitions and theorems to be reused across different mathematical domains:

```
class TopologicalSpace (α : Type*) where
  IsOpen : Set α → Prop
  isOpen_univ : IsOpen univ
  isOpen_inter : ∀ s t, IsOpen s → IsOpen t → IsOpen (s ∩ t)
  isOpen_sUnion : ∀ s, (∀ t ∈ s, IsOpen t) → IsOpen (sUnion s)
```

Key topological concepts used in our formalization include connectedness (formalized as `IsConnected`), path-connectedness (using `IsPathConnected` and the unit interval type `unitInterval`), continuous functions (`Continuous f`, with local variants `ContinuousAt` and `ContinuousOn`), and closure operations (`closure : Set α → Set α`). Lean treats sets as predicates (`Set α := α → Prop`), where set membership is simply function application.

4 Formalizing the topologist’s sine curve

As part of my thesis work, with the help and revision from Prof David Loeffler, I have formalized a well-known counterexample in topology: the **topologist’s sine curve**. This classic example illustrates a space that is **connected** but not **path-connected**. My original proof follows Conrad’s paper ([Con]), with a few modifications and some differences from the final formalization **Counterexamples – Topologist’s Sine Curve**. The topologist’s sine curve is defined as the graph of $y = \sin(1/x)$ for $x \in (0, \infty)$, together with the origin $(0, 0)$. We define three sets in \mathbb{R}^2 :

- S : the oscillating curve $\{(x, \sin(1/x)) : x > 0\}$
- Z : the singleton set $\{(0, 0)\}$
- T : their union $S \cup Z$

In Lean, this is expressed as follows:

```
open Real Set
def pos_real := Ioi (0 : ℝ)
noncomputable def sine_curve := fun x ↦ (x, sin (x-1))
def S : Set (ℝ × ℝ) := sine_curve '' pos_real
def Z : Set (ℝ × ℝ) := { (0, 0) }
def T : Set (ℝ × ℝ) := S ∪ Z
```

We open the `Real` and `Set` namespaces to avoid prefixing real number and set operations with `Real.` and `Set.`, respectively. We define the interval $(0, \infty)$ as `pos_real`, using the predefined notation `Ioi 0`, from `Set`. The function `sine_curve` maps a positive real number to a point on the topologist’s sine curve in \mathbb{R}^2 . Here, `''` denotes the image of a set under a function. It’s noncomputable because it involves the sine function, which is not computable in Lean’s core logic. The sets `S`, `Z`, and `T` are defined using set operations, and `{ (0, 0) }`

denotes the singleton set containing the point $(0,0)$. `Set` is the type of sets, defined as predicates (i.e., functions from a type to `Prop`). The sets are subsets of the product space \mathbb{R}^2 , represented as $\mathbb{R} \times \mathbb{R}$. The sin function `sin` is defined in the `Real`.

The goal is to prove that T is connected but not path-connected. Let's start with connectedness.

4.1 T is connected

First of all one can directly see that S is connected, since it is the image of the set $((0,\infty))$ under the continuous map $x \mapsto (x, \sin(1/x))$ and a interval in \mathbb{R} is connected. Moreover, the closure of S is connected, and every set in between a connected set and its closure are connected. Since T is contained in the closure of S , T is connected. This is how a mathematician would argue informally, using known facts. However, in a formal proof, one must justify each step. For instance, justifying that S is connected requires proving that the map $x \mapsto (x, \sin(1/x))$ is continuous on $(0, \infty)$ and that $(0, \infty)$ is connected.

As we have seen, even showing that a rational number is non-negative requires several steps and the use of various lemmas from `Mathlib`. Similarly, proving that a set is connected can involve multiple steps for the newer programmer.

We can use the structure `IsConnected` to set up the statement and see if we can argue similarly in Lean.

```
lemma S_is_conn : IsConnected S := by sorry
```

In the file where `IsConnected` is defined, `Topology/Connected/Basic.lean`, we see that it requires S to be nonempty and preconnected. You can verify this by unfolding `IsConnected` in the goal.

```

lemma S_is_conn : IsConnected S := by
  unfold IsConnected
  ⊢ S.Nonempty ∧ IsPreconnected S
  sorry

```

Following the definition of `IsPreconnected`, we see that it captures the usual definition of preconnectedness: that S cannot be partitioned into two nonempty disjoint open sets. This trivially requires nonemptiness to make sense. The `unfold` tactic helps to expand definitions; one can use it to expand the definition of S or `pos_real` defined before, as well as other Mathlib expressions. Reflecting our argument, we can check if Mathlib includes the fact that every interval is connected and that connectedness is preserved under continuous maps. Indeed, in `Topology/Connected/Interval.lean`, we find the theorem `isConnected_Ioi.image`, stating that the image of an interval of the form (a, ∞) under a continuous map is connected.

```

lemma S_is_conn : IsConnected S := by
  apply isConnected_Ioi.image
  -- ⊢ ContinuousOn sine_curve (Ioi 0)
  sorry

```

The `apply` tactic applies the theorem similar to `exact`, the latter tries to close the goal with `rf1`. The theorem `isConnected_Ioi.image` requires proving the continuity of the map on the interval $(0, \infty)$, which is expressed as `ContinuousOn sine_curve (Ioi 0)`. The predicate `ContinuousOn f S` expresses that a function f is continuous on a set S , which is what we need to prove now. The function $x \mapsto (x, \sin(1/x))$ is continuous on $(0, \infty)$ as the product of two functions continuous on the given domain: the identity map $x \mapsto x$ and the map $x \mapsto \sin(1/x)$.

Here is the full proof in Lean:

```

lemma inv_is_continuous_on_pos_real : ContinuousOn (fun x : ℝ => x-1)
  (pos_real) := by
  apply ContinuousOn.inv₀
  · exact continuous_id.continuousOn
  · intro x hx; exact ne_of_gt hx

lemma sin_comp_inv_is_continuous_on_pos_real : ContinuousOn
  (sine_curve) (pos_real) := by
  apply ContinuousOn.prodMk continuous_id.continuousOn
  apply continuous_sin.comp_continuousOn
  exact inv_is_continuous_on_pos_real

```

Starting from the bottom lemma, `ContinuousOn.prodMk` states that the product of two functions continuous on a set is continuous on that set, requiring a proof of the continuity of each component. The first component is the identity map, which is continuous on any set. Mathlib provides `continuous_id.continuousOn` for this purpose. The second component is the composition of the sine function with the inverse function. The sine function is continuous everywhere, and for this we can use `continuous_sin`. The method `comp_continuousOn` is accessible from the fact that `continuous_sin` gives an instance of a continuous map and is generalized in the `ContinuousOn` module. The theorem `Continuous.comp_continuousOn` states that the composition of a continuous function with a function that is continuous on a set is continuous on that set, and requires proof of the continuity on the set of the inner function. We separate the proof that the inverse function is continuous on the positive reals into the auxiliary lemma `inv_is_continuous_on_pos_real`. The theorem `continuousOn.inv₀` states that if a function is continuous and non-zero on a set, then its inverse is continuous on that set. The continuity of the identity map is proved as before. The second argument requires proving that $x \neq 0$ for all x in $(0, \infty)$.

```

· intro x hx
  exact ne_of_gt hx

```

The hypothesis `hx` states that x is in $(0, \infty)$, which implies that $x > 0$. The theorem `ne_of_gt` states that if a real number is greater than zero, then it is non-zero, which completes the proof. Thus the final proof goes as follows:

```

lemma S_is_conn : IsConnected S := by
  apply isConnected_Ioi.image
  · exact sin_comp_inv_is_continuous_on_pos_real

```

When writing a proof, one starts by working out the informal argument on paper. Then one tries to translate it into Lean, step by step, looking for theorems in Mathlib. Afterwards, one can try to optimize the proof by removing unnecessary steps or refactoring it. Proving properties like continuity and connectedness is very common, and there are obviously ways to achieve this with less work. Let's showcase a refactoring of the entire proof. First, the auxiliary lemmas can be reduced to one-liners.

```
lemma inv_is_continuous_on_pos_real : ContinuousOn (fun x : ℝ => x-1)
  (pos_real) :=
  ContinuousOn.inv₀ (continuous_id.continuousOn) (fun _ hx => ne_of_gt
    hx)
```

```
lemma sin_comp_inv_is_continuous_on_pos_real : ContinuousOn
  (sine_curve) (pos_real) :=
  ContinuousOn.prodMk continuous_id.continuousOn <|
    Real.continuous_sin.comp_continuousOn <|
      (inv_is_continuous_on_pos_real)
```

We removed the `by` keyword since we can provide a **term** that directly proves the statement. In `inv_is_continuous_on_pos_real`, we directly apply `ContinuousOn.inv₀` with the two required arguments. Notice that we can use a lambda function `fun _ hx => ne_of_gt hx` to prove that $x \neq 0$ for all x in $(0, \infty)$ (recall the propositions-as-types correspondence). In the next lemma, we use the `<|` reverse application operator, which allows us to avoid parentheses by changing the order of application. This means that `f < g <| h` is equivalent to `f (g h)`. We can inline these two lemmas into the main proof to get a final one-liner:

```
lemma S_is_conn : IsConnected S :=
  isConnected_Ioi.image sine_curve <| continuous_id.continuousOn.prodMk
    <|
      continuous_sin.comp_continuousOn <|
        ContinuousOn.inv₀ continuous_id.continuousOn (fun _ hx => ne_of_gt
          hx)
```

Notice again the use of the **pipe** operator. Reading from left to right, we are building up the proof by successive applications:

- We start with `isConnected_Ioi.image sine_curve`, which states that the image of $(0, \infty)$ under `sine_curve` is connected if we can prove the function is continuous.
- We then apply `continuous_id.continuousOn.prodMk`, which constructs the product of two continuous functions.
- Next, `continuous_sin.comp_continuousOn` provides the continuity of the sine composition.
- Finally, `ContinuousOn.inv₀ continuous_id.continuousOn (fun _ hx => ne_of_gt hx)` proves the continuity of the inverse function on positive reals.

The entire chain can be read as building the continuity proof from the innermost function (the inverse) outward to the complete sine curve function, which is then used to prove that S is connected.

Since the intersection of Z and S is empty, we cannot directly conclude that T is connected from the connectedness of its components alone. However, we can use the fact that every subset between a connected set and its closure is connected.

Theorem 4.1 *Let C be a connected topological space, and denote \overline{C} as its closure. It follows that every subset $C \subseteq S \subseteq \overline{C}$ is connected.*

In Mathlib, this theorem is available as `IsConnected.subset_closure`. We can set up the statement and progress from there.

```
theorem T_is_conn : IsConnected T := by
  apply IsConnected.subset_closure
  · exact S_is_conn --  $\vdash \text{IsConnected } ?s$ 
  · tauto_set --  $\vdash S \subseteq T$ 
  · sorry --  $\vdash T \subseteq \text{closure } S$ 
```

The theorem requires three goals:

1. That S is connected, which was already proved in `S_is_conn`.
2. That $S \subseteq T$, which is a trivial set operation. The tactic `tauto_set` handles this kind of set tautologies.
3. That $T \subseteq \overline{S}$ (the closure of S), which requires proof.

Let's continue with the final point.

```
lemma T_sub_cls_S : T  $\subseteq$  closure S := by
  intro x hx
  cases hx with
  | inl hxS => exact subset_closure hxS
  | inr hxZ =>
    sorry
```

Proving that one set is contained in another can be done naively in a pointwise manner. We introduce an element $x \in \mathbb{R}^2$ together with the proof that $x \in T$. Since T is a union, we use `cases` to separate the two cases. When $x \in S$, the goal is trivially solved by `exact subset_closure hxS`. The case where $x \in Z$, requires more work.

Now a trick. One of the most painful issue in formalizing math in Lean is the use of existing theorems. One can use several ways to look for the exact theorem. Let's try using the `apply?` tactic to see what the infoview suggests:

```
lemma T_sub_cls_S : T  $\subseteq$  closure S := by
  intro x hx
  cases hx with
  | inl hxS => exact subset_closure hxS
  | inr hxZ =>
    apply?
    sorry
```

Depending on the previous work in the file, Lean can already unify the goal with available theorems and suggest the next step. Similar tactics include:

- `exact?` for finding an exact match to close the goal
- `rw?` for suggesting rewrites and definitionally equal replacements

- `simp?` for suggesting simplifications

Another useful tool is Loogle (similar to Haskell’s Hoogle), which helps you find theorems by their type signature or name patterns. You can access it at <https://loogle.lean-lang.org/> or use it directly in VS Code. The best approach is, anyway, to think first about how you would tackle the problem on a piece of paper, as mentioned earlier. Since we are working with a topology on \mathbb{R} , we know that this is a metrizable topology, therefore it is induced by the metric space structure on \mathbb{R} . Thus, we can expand our toolkit by working with Lean’s `MetricSpace` module, which provides specialized tools for reasoning about metric spaces, such as balls, distances, and metric-specific characterizations of continuity and convergence. We know that the closure of a set contains all its limit points. To show that the point $(0, 0)$ is contained in the closure of S , we need to show that it is a limit point of S . Thus, one can define a sequence in S tending to $(0, 0)$, and we are done. At this point, `apply?` suggests several ways to proceed, involving new symbols such as \mathcal{N} (neighborhoods) and f (eventually/frequently). For example:

```
Try this: refine Frequently.mem_closure ?_
Remaining subgoals:
  ⊢ ∃f (x : ℝ × ℝ) in ℳ x, x ∈ S
```

or the more familiar metric space approach:

```
Try this: refine Metric.mem_closure_iff.mpr ?_
Remaining subgoals:
  ⊢ ∀ ε > 0, ∃ b ∈ S, dist x b < ε
```

While we could work with metric space properties directly using the familiar ε - δ formulation, we instead introduce and explain filters. This approach may seem more abstract initially, but it provides a more general and powerful framework that works uniformly across all topological spaces, not just metric spaces. Moreover, once understood, filters often make proofs more concise and elegant by allowing us to reason at a higher level of abstraction.

4.1.1 Limits and Convergence with Filters

As seen before, continuity is defined in the `Topology` type class. One can create their own epsilon-delta definition given a notion of distance. Similarly, defining the convergence of a sequence in \mathbb{R} can be done manually as follows:

```
def ConvergesTo (ε : ℝ > 0) (s : ℕ → ℝ) (a : ℝ) :=
  ∀ ε > 0, ∃ N, ∀ n ≥ N, |s n - a| < ε
```

Where s is a sequence and a is the limit point. However, there are many other types of limits to consider, like limits of a function at a point, limits at infinity (from above or below), one-sided limits (from the left or right) and so on. Defining each of these separately would require a huge amount of work to include in Mathlib, with significant duplication of theorems and proofs. Moreover, many fundamental theorems (like the characterization of continuity via limits)

would need to be reproved for each type of limit. Fortunately, Bourbaki solved this issue by introducing the notion of **filters** to unify all concepts of limits, convergence, neighborhoods and terms like eventually or frequently often into a single framework. Mathlib adopts this notion to achieve an elegant solution that fully covers the entire landscape of limit-related concepts. Intuitively, a filter represents a notion of "sufficiently large" subsets. More formally, a filter F on a type X is a collection of subsets of X satisfying three axioms:

1. **Non-emptiness:** $X \in F$ (the whole space is in the filter)
2. **Upward closure:** If $U \in F$ and $U \subseteq V$, then $V \in F$ (supersets of "large" sets are "large")
3. **Intersection closure:** If $U, V \in F$, then $U \cap V \in F$ (finite intersections of "large" sets are "large")

We are going to use some of the following concepts:

- **At top filter** `atTop` : `Filter N`: Contains sets that include all sufficiently large natural numbers. Formally, $U \in \text{atTop}$ if and only if there exists N such that $\{n \mid n \geq N\} \subseteq U$. This captures the idea of " $n \rightarrow \infty$."
- **Neighborhood filter** $\mathcal{N} x$: In a topological space, this filter contains all neighborhoods of the point x . A set is in $\mathcal{N} x$ if it contains an open set containing x . This captures the idea of "near x ."
- $\forall^f x \text{ in } f, p x (f.\text{Eventually } p)$: "Eventually in filter f , property p holds." This means there exists some set $U \in f$ such that p holds for all $x \in U$. For example, $\forall^f n \text{ in } \text{atTop}, n > 100$ means "for all sufficiently large n , we have $n > 100$."
- $\exists^f x \text{ in } f, p x (f.\text{Frequently } p)$: "Frequently in filter f , property p holds." This means for every set $U \in f$, there exists some $x \in U$ where p holds. This captures the idea that p holds "infinitely often" or "arbitrarily close." For example, $\exists^f n \text{ in } \text{atTop}, \text{Even } n$ means "there are arbitrarily large even numbers."
- **Tendsto** `f l1 l2`: "Function f tends from filter l_1 to filter l_2 ." This is used for convergence.

Example 4.2 Here some examples of the use of filter in Lean.

- *Sequence convergence:* `Tendsto s atTop (N a)` means $s_n \rightarrow a$ as $n \rightarrow \infty$
- *Function limits:* `Tendsto f (N x) (N y)` means $f(x') \rightarrow y$ as $x' \rightarrow x$
- *Continuity:* A function is continuous at x iff `Tendsto f (N x) (N (f x))` (does this remember the continuity in terms of open sets?)

Using filters, we can prove that $T \subseteq \overline{S}$ by showing that the origin is a limit point of S . We construct a sequence $f : \mathbb{N} \rightarrow \mathbb{R}^2$ in S converging to $(0, 0)$ using the `Tendsto` framework:

```

lemma T_sub_cls_S : T ⊆ closure S := by
  intro x hx
  cases hx with
  | inl hxS => exact subset_closure hxS
  | inr hxZ =>
    rw [hxZ]
    -- Define sequence: f(n) = (1/(nπ), 0)
    let f : ℕ → ℝ × ℝ := fun n => ((n * Real.pi)⁻¹, 0)
    -- Show f converges to (0, 0)
    have hf : Tendsto f atTop (ℕ (0, 0)) := by
      refine Tendsto.prodMk_nhds ?_ tendsto_const_nhds
      exact tendsto_inv_atTop_zero.comp
        (Tendsto.atTop_mul_const' Real.pi_pos
         tendsto_natCast_atTop_atTop)
    -- Show f eventually takes values in S
    have hf' : ∀f n in atTop, f n ∈ S := by
      filter_upwards [eventually_gt_atTop 0] with n hn
      exact ⟨(n * Real.pi)⁻¹,
        inv_pos.mpr (mul_pos (Nat.cast_pos.mpr hn) Real.pi_pos),
        by simp [f, sine_curve, inv_inv, Real.sin_nat_mul_pi]⟩
    -- Apply sequential characterization of closure
    exact mem_closure_of_tendsto hf hf'

```

The proof is already reduced as much as possible. Let's break down what's happening in without getting into details. Using **let**, we define $f(n) = (\frac{1}{n\pi}, 0)$, which we will show converges to $(0, 0)$ and stays in S .

1. **Convergence proof (hf)**: We show $\text{Tendsto } f \text{ atTop } (\mathbb{N} (0, 0))$.
 - We use `Tendsto.prodMk_nhds` to split the product: we need to show the first coordinate tends to 0 and the second is constantly 0.
 - For the first coordinate, we compose `tendsto_inv_atTop_zero` (which states $\frac{1}{x} \rightarrow 0$ as $x \rightarrow \infty$) with the fact that $n\pi \rightarrow \infty$.
 - The second constant coordinate is handled by `tendsto_const_nhds`.
2. **Membership proof (hf')**: We show $\forall^f n \text{ in atTop}, f n \in S$, meaning $f(n) \in S$ for all sufficiently large n .
 - We use `filter_upwards`, which allows us to combine hypotheses about properties that hold eventually to prove another property holds eventually. Here, we combine it with `eventually_gt_atTop 0`, which states that eventually $n > 0$.
 - For such n , we show $f(n) = (\frac{1}{n\pi}, 0)$ is in S by noting that the second term is:

$$\sin\left(\frac{1}{\left(\frac{1}{n\pi}\right)}\right) = \sin(n\pi) = 0.$$

Finally, `mem_closure_of_tendsto` combines these facts: if a sequence eventually stays in S and converges to x , then x is in the closure of S .

4.1.2 Finalising the first part of the proof

If you are a one-liner enthusiast like me, you don't mind trying to combine bits and pieces to get a clean final result. We can simplify the final theorem as follows initially:

```
theorem T_is_conn : IsConnected T :=
  IsConnected.subset_closure S_is_conn (by tauto_set) T_sub_cls_S
```

The second argument is still in tactic mode with `by tauto_set`, but it looks clean and we can keep it as is. With a bit of courage, we can also inline the proof of `S_is_conn` (while `T_sub_cls_S` is way too long to inline) to get a more self-contained one-liner:

```
theorem T_is_conn : IsConnected T :=
  IsConnected.subset_closure (isConnected_Ioi.image sine_curve <|
    continuous_id.continuousOn.prodMk <|
    Real.continuous_sin.comp_continuousOn <|
    ContinuousOn.inv₀ continuous_id.continuousOn
    (fun _ hx => ne_of_gt hx)) (by tauto_set) T_sub_cls_S
```

Making these amendments is not only for the sake of shortening the proof. Lean will, obviously, compile the proof faster by not entering tactic mode or using multiple tactics. Tactics internally hide many operations they automatically perform to close the goal. Moreover, if we directly provide a term for the proof, Lean will infer and unify everything by definitional equality; remember the very first example we did. By providing explicit proof terms, we give Lean less work to do, making the proof more transparent and efficient. This practice of "golfing" is essential in a huge library such as Mathlib community that needs to balance performance and maintainability. From now on the rest of the code will be presented in its reduced form. Here is the link to the entire first part of the proof: [\[link to Lean live\]](#)

Note 4.3 *The proof merged into the Mathlib library takes Z as $\{0\} \times [-1, 1]$ instead of the singleton $\{(0, 0)\}$. This, together with the fact that T equals the closure of S , yields a stronger and more general result. The Mathlib version demonstrates that the entire vertical segment at $x = 0$ lies in the closure of the oscillating curve, providing a more complete characterization of the topologist's sine curve. This stronger version shows that a closed set (specifically, the closure of S) can be connected but not path-connected. Showing that forming closure can destroy the property of path connectedness for subsets of a topological space.*

4.2 T is not path-connected

The main and most substantial part is showing that T is not path-connected. Showing this informally already requires constructing and pointing out various steps in order to convince an ideal reader. One can argue informally by contradiction: suppose a path exists in the topologist's sine curve T connecting a point in S to a point in Z . As the path approaches the y -axis (where $x \rightarrow 0$),

the y -coordinate must oscillate infinitely between -1 and 1 due to the behavior of $\sin(1/x)$ as $x \rightarrow 0^+$. This infinite oscillation contradicts the continuity of the path, which is a fundamental requirement for path-connectedness. To be more precise, we need to construct a sequence that it eventually oscillates, establishing the contradiction. We start by setting up the theorem:

```
theorem T_is_not_path_conn : ¬ (IsPathConnected T) :=
  by sorry
```

In mathematics, we normally define a path-connected space as follows.

Definition 4.4 *A topological space X is said to be path-connected if for every two points $a, b \in X$, there exists a path, i.e., a continuous map $p : [0, 1] \rightarrow X$ such that $p(0) = a$ and $p(1) = b$.*

The interval $[0, 1]$ is the standard choice for the domain of paths. In Mathlib, `PathConnectedSpace X` is a type class that asserts the entire topological space X is path-connected, while `IsPathConnected S` is a predicate used to infer that a subset S of a topological space is path-connected.

```
def IsPathConnected (F : Set X) : Prop :=
  ∃ x ∈ F, ∀ {y}, y ∈ F → JoinedIn F x y
```

The auxiliary predicate `JoinedIn` is defined as:

```
def JoinedIn (S : Set X) (x y : X) : Prop :=
  ∃ γ : Path x y, ∀ t, γ t ∈ S
```

where `Path x y` denotes a continuous map $\gamma : [0, 1] \rightarrow X$ with $\gamma(0) = x$ and $\gamma(1) = y$. Mathlib uses `unitInterval` as the standard definition for $[0, 1]$ in constructions such as the definition of a path.

Now let's start with the first part of the proof:

```

theorem T_is_not_path_conn : ¬ (IsPathConnected T) := by
  -- Assume we have a path from z = (0, 0) to w = (1, sin(1))
  have hz : z ∈ T := Or.inr rfl
  have hw : w ∈ T := Or.inl ⟨1, ⟨zero_lt_one' ℝ, rfl⟩⟩
  intro p_conn
  apply IsPathConnected.joinedIn at p_conn
  specialize p_conn z hz w hw
  let p := JoinedIn.somePath p_conn

```

We introduce two points: $z = (0,0)$ and $w = (1, \sin(1))$, and prove they are both in T . Using `intro p_conn`, we assume that T is path-connected. Notice that the goal is now `False`, meaning we must find a contradiction. The last three lines extract an explicit path p connecting z and w :

- `apply IsPathConnected.joinedIn at p_conn` transforms the path-connectedness assumption into the statement that any two points in T are joined.
- `specialize p_conn z hz w hw` specializes this to our specific points z and w .
- `let p := JoinedIn.somePath p_conn` extracts a concrete path from the existential statement.

Conrad's paper ([Con]) defines a time $t_0 \in [0, 1]$ as the first time the path p jumps from $(0,0)$ to the graph of $\sin(1/x)$, where the x-coordinate map $(x : \mathbb{R}^2 \rightarrow \mathbb{R})$ of p is positive.

$$t_0 = \inf\{t \in [0, 1] : x(p(t)) > 0\}$$

The argument then uses the continuity of the x -coordinate map composed with the path p . By continuity at t_0 , we can find a neighborhood around t_0 where the path stays close to $(0,0)$. Specifically, with $\varepsilon = 1/2$, there exists $\delta > 0$ such that for all t with $|t - t_0| < \delta$, we have $\|p(t) - p(t_0)\| < 1/2$. We want to show the oscillating behavior around $(0,0)$ indeed. To simplify some steps, we instead define

$$t_0 = \sup\{t \in [0, 1] : x(p(t)) = 0\}$$

to be the last time the path remains at $(0,0)$. The same continuity argument applies with this definition.

```

-- Consider the composition of the x-coordinate map with p, which is
  continuous
have xcoord_pathcont : Continuous fun t ↦ (p t).1 :=
  continuous_fst.comp p.continuous
-- Let t₀ be the last time the path is on the y-axis
let t₀ : unitInterval := sSup {t | (p t).1 = 0}
let xcoord_path := fun t => (p t).1
-- The x-coordinate of the path at t₀ is 0
have hpt₀_x : (p t₀).1 = 0 :=
  (isClosed_singleton.preimage xcoord_pathcont).sSup_mem ⟨0, by aesop⟩
-- By continuity of the path, we can find a δ > 0 such that
-- for all t in [t₀ - δ, t₀ + δ], ||p(t) - p(t₀)|| < 1/2
-- Hence the path stays in a ball of radius 1/2 around (0, 0)
obtain ⟨δ, hδ, ht⟩ : ∃ δ > 0, ∀ t, dist t t₀ < δ →
  dist (p t) (p t₀) < 1/2 :=
  Metric.eventually_nhds_iff.mp <| Metric.tendsto_nhds.mp
  (p.continuousAt t₀) _ one_half_pos

```

The final statement uses the `obtain` tactic to extract witnesses from an existential statement. This tactic deconstructs the existential quantifier $\exists \delta > 0, \dots$ into concrete values: δ (the distance), $h\delta$ (the proof that $\delta > 0$), and ht (the proof that the distance condition holds). Since \mathbb{R}^2 is a metric space, we can work with the distance function $\text{dist} : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$, which computes the Euclidean distance between two points. The statement $\text{dist } t \ t_0 < \delta$ expresses $|t - t_0| < \delta$ in the unit interval, while $\text{dist } (p \ t) \ (p \ t_0) < 1/2$ expresses $\|p(t) - p(t_0)\| < 1/2$ in \mathbb{R}^2 . The proof itself leverages `Metric` module.

- `p.continuousAt t₀` asserts that the path p is continuous at t_0
- `Metric.tendsto_nhds.mp` converts this to the metric space characterization: for any $\varepsilon > 0$, there exists $\delta > 0$ such that points within δ of t_0 map to points within ε of $p(t_0)$
- `Metric.eventually_nhds_iff.mp` further unpacks this into the $\forall t, \text{dist } t \ t_0 < \delta \rightarrow \text{dist } (p \ t) \ (p \ t_0) < \varepsilon$ form
- We instantiate with $\varepsilon = 1/2$ using `one_half_pos`

We can find a time t_1 greater than t_0 that remains in the neighborhood of t_0 , and obtain a point $a = x(p(t_1)) > 0$ which is positive.

```

-- Let  $t_1$  be a time when the path is not on the y-axis
--  $t_1$  is in  $(t_0, t_0 + \delta]$ , hence  $t_1 > t_0$ 
obtain ⟨ $t_1$ ,  $ht_1$ ⟩ :  $\exists t_1, t_1 > t_0 \wedge \text{dist } t_0 \ t_1 < \delta$  := by
  let  $s_0$  := ( $t_0$  :  $\mathbb{R}$ ) -- cast  $t_0$  from unitInterval to  $\mathbb{R}$  for manipulation
  let  $s_1$  := min ( $s_0 + \delta/2$ ) 1
  have  $hs_0\_delta\_pos$  :  $0 \leq s_0 + \delta/2$  := add_nonneg  $t_0$ .2.1 (by positivity)
  have  $hs_1$  :  $0 \leq s_1$  := le_min  $hs_0\_delta\_pos$  zero_le_one
  have  $hs_1'$  :  $s_1 \leq 1$  := min_le_right ..
  sorry
-- Let  $a = xcoord\_path \ t_1 > 0$ 
-- This follows from the definition of  $t_0$  and  $t_0 < t_1$ 
-- so  $t_1$  must be in  $S$ , which has positive x-coordinate
let  $a$  := ( $p \ t_1$ ).1
have  $ha$  :  $a > 0$  := by
  obtain ⟨ $x$ ,  $hxI$ ,  $hx\_eq$ ⟩ :  $p \ t_1 \in S$  := by
    cases  $p\_conn.somePath\_mem \ t_1$  with
    | inl  $hS$  => exact  $hS$ 
    | inr  $hZ$  =>
      -- If  $p \ t_1 \in Z$ , then  $(p \ t_1).1 = 0$ 
      have :  $(p \ t_1).1 = 0$  := by rw [ $hZ$ ]
      -- So  $t_1 \leq t_0$ , contradicting  $t_1 > t_0$ 
      have  $hle$  :  $t_1 \leq t_0$  := le_sSup this
      have  $hle\_real$  :  $(t_1 : \mathbb{R}) \leq (t_0 : \mathbb{R})$  := Subtype.coe_le_coe.mpr  $hle$ 
      have  $hgt\_real$  :  $(t_1 : \mathbb{R}) > (t_0 : \mathbb{R})$  := Subtype.coe_lt_coe.mpr
       $ht_1.1$ 
      linarith
  simp only [ $a$ ,  $\leftarrow hx\_eq$ ] using  $hxI$ 

```

The code is quite convoluted in Lean, and i will omit a detailed explanation as well as some part of it. However, it's worth mentioning a few key technical points. The type `unitInterval` is a **subtype** of \mathbb{R} , defined as $\{x : \mathbb{R} \mid 0 \leq x \leq 1\}$. In Lean, a subtype $\{x : \alpha \mid P \ x\}$ bundles a value x of type α together with a proof that x satisfies the predicate P . Manipulating terms of `unitInterval` directly is challenging because this type lacks many algebraic operations such as addition, minimum, etc. Therefore, we cast to \mathbb{R} (with `let s_0 := (t_0 : \mathbb{R})`) to perform arithmetic operations, then cast back to `unitInterval` by providing proofs that the bounds $[0, 1]$ are satisfied (hs_1, hs_1'). In the second case of the inner statment of `have ha : $a > 0$` , if $p(t_1) \in Z$, then $(p \ t_1).1 = 0$ by definition of $Z = \{(0, 0)\}$. This implies $t_1 \leq t_0$ by the definition of t_0 as the supremum. However, we also have $t_1 > t_0$ from our construction of t_1 ($ht_1.1$). The tactic `linarith`, an automated solver for linear arithmetic, recognizes this contradiction by observing both $hle_real : (t_1 : \mathbb{R}) \leq (t_0 : \mathbb{R})$ and $hgt_real : (t_1 : \mathbb{R}) > (t_0 : \mathbb{R})$. Since these statements are contradictory, `linarith` proves `False`. Lemmas like `Subtype.coe_lt_coe` allow us to transfer inequalities between the subtype and its underlying type, needed for `linarith`.

Finally, `simp only [a , $\leftarrow hx_eq$] using hxI` completes the proof. The tactic `simp` combines simplification (`simp`) with assumption matching. The directive `only [a , $\leftarrow hx_eq$]` unfolds the definition of $a = (p \ t_1).1$ and rewrites

using `hx_eq` in the reverse direction, transforming the goal from $(p \ t_1).1 > 0$ to $(\text{sine_curve } x).1 > 0$. Since $\text{sine_curve } x = (x, \sin(1/x))$, this simplifies to $x > 0$, which is exactly the hypothesis `hxI`. The `using hxI` clause applies this hypothesis to close the goal.

Next, the image $x(p([t_0, t_1]))$ is connected (as the continuous image of a connected set), and it contains $0 = x(p(t_0))$ and $a = x(p(t_1))$. Since every connected subset of \mathbb{R} is an interval, we have

$$[0, a] \subseteq x(p([t_0, t_1]))$$

This will be crucial for the next step, where we show that the path must oscillate.

```
-- The image x(p([t0, t1])) is connected and contains 0 and a
-- Therefore [0, a] ⊆ x(p([t0, t1]))
have Icc_of_a_b_sub_Icc_t0_t1: Set.Icc 0 a ⊆ xcoord_path '' Set.Icc
  t0 t1 :=
  IsConnected.Icc_subset
    ((isConnected_Icc (le_of_lt ht1.1)).image _
  xcoord_pathcont.continuousOn)
    ((t0, left_mem_Icc.mpr (le_of_lt ht1.1), hpt0_x))
    ((t1, right_mem_Icc.mpr (le_of_lt ht1.1), rfl))
```

Now we construct a sequence that demonstrates the contradiction. Recall that $\sin(\theta) = 1$ if and only if $\theta = \frac{(4k+1)\pi}{2}$ for some $k \in \mathbb{Z}$. Therefore, $(x, \sin(1/x)) = (x, 1)$ when

$$x = \frac{2}{(4k+1)\pi}$$

for $k \in \mathbb{N}$. As $k \rightarrow \infty$, these x -values approach 0, so infinitely many of them lie in any interval $[0, a]$. We define this sequence and establish its key properties:

```
noncomputable def xs_pos_peak := fun (k : ℕ) => 2 / ((4 * k + 1) *
  Real.pi)
lemma xs_pos_peak_tendsto_zero : Tendsto xs_pos_peak atTop (ℕ 0) :=
  sorry
lemma xs_pos_peak_nonneg : ∀ k : ℕ, 0 ≤ xs_pos_peak k := sorry
lemma sin_xs_pos_peak_eq_one (k : ℕ) : Real.sin ((xs_pos_peak k)-1) = 1
  := sorry
```

The crucial property is that this sequence eventually enters $[0, a]$:

```
-- For any k ∈ ℕ, sin(1/xs_pos_peak(k)) = 1
-- Since xs_pos_peak converges to 0 as k → ∞,
-- there exist indices i ≥ 1 for which xs_pos_peak i ∈ [0, a]
have xpos_has_terms_in_Icc_of_a_b : ∃ i : ℕ, i ≥ 1 ∧ xs_pos_peak i ∈
  Set.Icc 0 a := sorry
```

This gives us points on the topologist's sine curve with y -coordinate equal to 1, lying arbitrarily close to the y -axis.

Now we can establish the final contradiction. Since $[0, a] \subseteq x(p([t_0, t_1]))$ by the previous argument, and $\text{xs_pos_peak}(i) \in [0, a]$ for some i , there must exist some $t' \in [t_0, t_1]$ such that $x(p(t')) = \text{xs_pos_peak}(i)$. This means

$p(t') = (\text{xs_pos_peak}(i), \sin(1/\text{xs_pos_peak}(i))) = (\text{xs_pos_peak}(i), 1)$, so the y -coordinate of $p(t')$ equals 1. However, since $t' \in [t_0, t_1] \subseteq [t_0, t_0 + \delta)$, we have $\text{dist}(t', t_0) < \delta$, which by our earlier continuity argument implies $\|p(t') - p(t_0)\| < 1/2$. But $\|p(t') - (0, 0)\| \geq |(p(t')).2| = |1| = 1 > 1/2$, yielding a contradiction.

```
-- Show there exists time t' in [t_0, t_1] ⊆ [t_0, t_0 + δ) such that p(t') =
  (*, 1)
obtain ⟨t', ht', hpath_t'⟩ : ∃ t' ∈ Set.Icc t_0 t_1, (p t').2 = 1 := sorry
-- Derive the final contradiction using t', ht', hpath_t'
-- First show that p t_0 = (0, 0)
have hpt_0 : p t_0 = (0, 0) := sorry
-- t' is within δ of t_0 (since t' ∈ [t_0, t_1] and dist t_0 t_1 < δ)
have t'_close : dist t' t_0 < δ := by
  calc dist t' t_0
    ≤ dist t_1 t_0 := dist_right_le_of_mem_uIcc (Icc_subset_uIcc' ht')
    _ = dist t_0 t_1 := dist_comm _ _
    _ < δ := ht_1.2
-- By continuity, p(t') should be close to p(t_0)
have close : dist (p t') (p t_0) < 1/2 := ht t' t'_close
-- But p(t') has y-coordinate 1, so it's actually far from p(t_0) = (0, 0)
have far : 1 ≤ dist (p t') (p t_0) := by
  calc 1 = |(p t').2 - (p t_0).2| := by simp [hpath_t', hpt_0]
    _ ≤ ‖p t' - p t_0‖ := norm_ge_abs_snd
    _ = dist (p t') (p t_0) := by rw [dist_eq_norm]
-- This is a contradiction: 1 ≤ dist (p t') (p t_0) < 1/2
linarith
```

4.3 T is connected not path-connected

Finally, we combine the two parts in the following concise and pleasant theorem:

```
theorem T_is_conn_not_pathconn : IsConnected T ∧ ¬IsPathConnected T :=
  ⟨T_is_conn, T_is_not_path_conn⟩
```

And now, since this code compiles successfully, these two lines stand as verified witnesses to the correctness of our entire proof. This showcases the power of proof assistants and formal reasoning: mathematics becomes not only more rigorous but also automatically verifiable. Furthermore, the formalization becomes a learning tool in its own right. Future readers can inspect each part of the code. Here the full proof: [\[link to Lean live\]](#)

References

- [Com20] The mathlib Community. “The Lean Mathematical Library”. In: *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs*. 2020.

- [Con] Keith Conrad. *Spaces that are connected but not path-connected*. <https://kconrad.math.uconn.edu/blurbs/topology/connnotpathconn.pdf>.
- [GTL89] Jean-Yves Girard, Paul Taylor, and Yves Lafont. *Proofs and Types*. Cambridge University Press, 1989.
- [Lea25a] Lean Mathematical Library community. *Mathematics in Lean*. https://leanprover-community.github.io/mathematics_in_lean/mathematics_in_lean.pdf. 2025.
- [Lea25b] Lean Mathematical Library community. *mathlib — The Lean Mathematical Library*. <https://github.com/leanprover-community/mathlib>. 2025.
- [LM20] Robert Y. Lewis and Paul-Nicolas Madelaine. “Simplifying Casts and Coercions”. In: *arXiv preprint arXiv:2001.10594* (2020).
- [NPS90] Bengt Nordström, Kent Petersson, and Jan M. Smith. *Programming in Martin-Löf’s Type Theory: An Introduction*. Oxford University Press, 1990.
- [Tho99] Simon Thompson. *Type Theory And Functional Programming*. University of Kent, 1999.
- [Wad15] Philip Wadler. “Propositions as Types”. In: *Communications of the ACM* 58.12 (2015), pp. 75–84.
- [WB89] Philip Wadler and Stephen Blott. “How to Make Ad-Hoc Polymorphism Less Ad Hoc”. In: *16th ACM Symposium on Principles of Programming Languages*. 1989.