

Contents

1	Introduction	1
1.1	Lean first steps	2
2	Propositional and Predicate Logic	4

1 Introduction

This serves as a brief starting point for understanding how mathematical proofs can be formalized in Lean, as well as being an introduction to the language itself. Lean is both a **functional programming language** and a **theorem prover**. We'll focus primarily on its role as a theorem prover. But what does this mean, and how can that be achieved?

A programming language defines a **set of rules, semantics, and syntax** for writing programs. To achieve a goal, a programmer must write a program that meets given specifications. There are two primary approaches: **program derivation** and **program verification** ([NPS90] Section 1.1). In **program verification**, the programmer first writes a program and then proves it meets the specifications. This approach checks for errors at **run-time** when the code executes. In **program derivation**, the programmer writes a proof that a program with certain properties exists, then extracts a program from that proof. This approach enables specification checking at **compilation-time**, catching errors while typing, thus, before execution. This distinction corresponds to **dynamic** versus **static type systems**. Most programming languages combine both approaches; providing basic types for annotation and compile-time checking, while leaving the remaining checks to be performed at runtime.

Example 1.1 *In a dynamically typed language, like JavaScript, variables can change type after they are created. For example, a variable defined as a number can later be reassigned to a string:*

```
let price = 100; // Javascript internal system recognize this
                variable as a number
price = "100";  // "100" transform the number 100 to a string. This
                is valid in JavaScript
```

TypeScript is a statically typed superset of JavaScript. Unlike JavaScript, it performs type checking at compile time. This means we can prevent the previous behavior, while writing our code, simply by adding a type annotation:

```
let price: number = 100;
price = "100"; // Error: Type 'string' is not assignable to type
               'number'
```

Now converting a variable with type annotation number to a string results into a compile time error.

Nonetheless, TypeScript, even though it has a sophisticated type system, cannot fully capture complex mathematical properties. As well as for the most programming languages, program specifications can only be enforced at runtime. Lean, by contrast, uses a much more powerful type system that enables it to express and verify mathematical statements with complete rigor fully during compilation time. This makes it particularly suitable as a **theorem prover** for formalizing mathematics.

Lean’s type system is based on **dependent type theory**, specifically the **Calculus of Inductive Constructions** (CIC) with various extensions. It’s important to note that **type theory** is not a single, unified theory, but rather a family of related theories with various extensions, ongoing developments, and rich historical ramifications.

Type theory emerged as an alternative foundation for mathematics, addressing paradoxes that arose in naive set theory. Consider Russell’s famous paradox: let $S = \{x \mid x \notin x\}$ be the set of all sets that do not contain themselves. This construction is paradoxical, leading to the contradiction $S \in S \iff S \notin S$. Type theory resolves such issues by working with **types** as primary objects rather than sets, and by restricting which constructions are well-formed.

Dependent type theory, the framework underlying Lean, extends basic type systems by allowing types to depend on values. For instance, one can define the type of vectors of length n , where n is itself a value. This capability makes dependent type theory particularly expressive for formalizing mathematics.

Various proof assistants have been developed based on different variants of type theory, including Agda, Coq, Idris, and Lean. Each system makes different design choices regarding which rules and features to include. Lean adopts the **Calculus of Inductive Constructions**, which extends the Calculus of Constructions, introduced in Coq, with **inductive types**. Inductive types allow for the definition of structures such as natural numbers, lists, and trees.

A fundamental design feature of Lean is its **universe hierarchy** of types, with **Prop** (the proposition type) as a distinguished universe at the bottom. The **Prop** universe exhibits two special properties: **impredicativity** and **proof irrelevance**. Proof irrelevance means that all proofs of the same proposition are considered equivalent. What matters is whether a proposition can be proven, not which specific proof is given. This separation between propositions (**Prop**) and data types (**Type**) was first introduced in N.G. de Bruijn’s **AUTOMATH system** (1967) ([Tho99]).

For our purposes, we do not need to delve deeply into the theoretical foundations; instead, we will introduce the relevant concepts as needed while working with Lean. The practical aspects of writing proofs in Lean will be our primary focus, and the theoretical machinery will be explained only insofar as it aids understanding of how to formalize mathematics effectively.

1.1 Lean first steps

In the language of type theory, and by extension in Lean, we write $x : X$ to mean that x is a **term** of type X . For example, $2 : \mathbb{N}$ annotates 2 as a natural

number, or more precisely, as a term of the natural number type. Lean has internally defined types such as `Nat` or `N` (you can type `\Nat` to get the Unicode symbol). The command `#check` allows us to inspect the type of any expression, term or variable.

Example 1.2

```
#check 2 -- 2 : Nat
#check 2 + 2 -- 2 + 2 : Nat
```

*[Try this example in Lean Web Editor] By following the link, you can try out the code in your browser. Lean provides a dedicated **infoview** panel on the right side. Position your cursor after `#check 2`, and the infoview will display the output `2 : Nat`. This dynamic interaction, where the infoview responds to your cursor position, is what makes Lean an **interactive theorem prover**. As you move through your code, the infoview continuously updates, showing computations, type information, and proof states at each location.*

At first glance, one might be tempted to view the colon notation as analogous to the membership symbol \in from set theory, treating types as if they were sets. While this intuition can be helpful initially, type theory offers a fundamentally richer perspective. The crucial insight is the **Curry-Howard correspondence**, also known as the **propositions-as-types** principle. This correspondence establishes a deep connection between mathematical proofs and programs: **propositions correspond to types**, and **proofs correspond to terms** inhabiting those types. Under this interpretation, a term $x : X$ can be understood in two more ways:

- As a **computational object**: x is a program or data structure of type X
- As a **logical object**: x is a proof of the proposition X

Lean is a concrete realization of the propositions-as-types principle, proving a theorem, within the language, amounts to constructing a term of the appropriate type. When we write `theorem_name : Proposition`, we are declaring that `theorem_name` is a proof (term) of `Proposition` (type). For example, consider proving that $2 + 2 = 4$:

Example 1.3

```
theorem two_plus_two_eq_four : 2 + 2 = 4 := rfl
```

Lean's syntax is designed to resemble the language of mathematics. Here, we use the `theorem` keyword to encapsulate our proof, of the statement/proposition $2 + 2 = 4$, giving it the name `two_plus_two_eq_four`. This allows us to reference and reuse this result later in our code. After the semicolon, `:`, we introduce the statement; $2 + 2 = 4$. The `:=` operator expects the proof term that establishes the theorem's validity. The proof itself consists of a single term: `rfl` (short for **reflexivity**). This is a proof term that works by **definitional equality**,

Lean’s kernel automatically reduces both sides of the equation to their normal (definitional) form and verifies they are identical. Since $2+2$ computes to 4, the proof succeeds immediately. We can now use this theorem in subsequent proofs. For instance:

```
example : 1 + 1 + 1 + 1 = 4 := two_plus_two_eq_four
```

Well, this example is simple enough for Lean to evaluate by itself: $1 + 1 + 1 + 1 = 2 + 2 = 4$ and conclude with `two_plus_two_eq_four`. Actually, `rfl` would solve the equation similarly, so this is just applying `rfl` again (it’s a bit of cheating). Here, I used `example`, which is handy for defining anonymous expressions for demonstration purposes. Before diving into the discussion, here is another keyword, `def`, used to introduce definitions and functions.

```
def addOne (n : Nat) : Nat := n + 1
```

This definition expects a natural number as its parameter, written `(n : Nat)` and returns a natural number. [Run in browser]

Let’s now turn to how logic is handled in Lean and how the Curry-Howard isomorphism is reflected concretely.

2 Propositional and Predicate Logic

Logic is the study of reasoning, branching into various systems. We refer to **classical logic** as the one that underpins much of traditional mathematics. It’s the logic of the ancient Greeks (not fair) and truth tables, and it remains used nowadays for pedagogical reasons. We first introduce **propositional logic**, which is the simplest form of classical logic. Later we will extend this to **predicate (or first-order) logic**, which includes quantifiers and predicates. In this setting, a **proposition** is a statement that is either true or false, and a **proof** is a logical argument that establishes the truth of a proposition. Propositions are constructed via **formulas** built from **propositional variables** (also called atomic propositions) combined with logical **connectives** such as “and” (\wedge), “or” (\vee), “not” (\neg), “implies” (\Rightarrow), and “if and only if” (\Leftrightarrow). These connectives allow the creation of complex or compound propositions.

Definition 2.1 (Propositional Formula) ([Tho99]) A *propositional formula* is either:

- A *propositional variable*: X_0, X_1, X_2, \dots , or
- A *compound formula* formed by combining formulas using connectives:

$$(A \wedge B), \quad (A \Rightarrow B), \quad (A \vee B), \quad \perp, \quad (A \Leftrightarrow B), \quad (\neg A)$$

where A and B are formulas themselves.

We are going to describe classical logic through a formal framework called **natural deduction system** developed by Gentzen in the 1930s ([Wad15]). It specifies rules for deriving **conclusions** from **premises** (assumptions from other propositions), called **inference rules**.

Example 2.2 (Deductive style rule) *Here is an hypothetical example of inference rule.*

$$\frac{P_1 \quad P_2 \quad \dots \quad P_n}{C}$$

Where the P_1, P_2, \dots, P_n , above the line, are hypothetical premises and, the hypothetical conclusion C is below the line.

The inference rules needed are:

- **Introduction rules** specify how to form compound propositions from simpler ones, and
- **Elimination rules** specify how to use compound propositions to derive information about their components.

Let's look at how we can define some connectives.

Conjunction (\wedge)

- Introduction

$$\frac{A \quad B}{A \wedge B} \wedge\text{-Intro}$$

- Elimination

$$\frac{A \wedge B}{A} \wedge\text{-Elim}_1$$

$$\frac{A \wedge B}{B} \wedge\text{-Elim}_2$$

Disjunction (\vee)

- Introduction

$$\frac{A}{A \vee B} \vee\text{-Intro}_1$$

$$\frac{B}{A \vee B} \vee\text{-Intro}_2$$

- Elimination (Proof by cases)

$$\frac{A \vee B \quad [A] \vdash C \quad [B] \vdash C}{C} \vee\text{-Elim}$$

Implication (\rightarrow)

- Introduction

$$\frac{[A] \vdash B}{A \rightarrow B} \rightarrow\text{-Intro}$$

- Elimination (Modus Ponens)

$$\frac{A \rightarrow B \quad A}{B} \rightarrow\text{-Elim}$$

Notation 2.3 We use $A \vdash B$ (called *turnstile*) to designate a deduction of B from A . It is employed in Gentzen's **sequent calculus** ([GTL89]), whereas in natural deduction the corresponding symbol is

$$\begin{array}{c} A \\ \vdots \\ B \end{array}$$

There are some minor differences, in fact, which I don't fully understand. The square brackets around a premise $[A]$ mean that the premise A is meant to be **discharged** at the conclusion. The classical example is the introduction rule for the implication connective. To prove an implication $A \rightarrow B$, we assume A (shown as $[A]$), derive B under this assumption, and then discharge the assumption A to conclude that $A \rightarrow B$ holds without the assumption. The turnstile is predominantly used in judgments and type theory with the meaning of "entails that".

Lean has its own syntax for connectives and their relative inference rules. For instance $A \wedge B$ can be presented as `And(A, B)` or $A \wedge B, .$ Its introduction rule is constructed by `And.intro _ _` or shortly $\langle _, _ \rangle$ (underscore are placeholder for assumptions or "propositional functions"). The pair $A \wedge B$ can be then consumed using elimination rules `And.left` and `And.right`.

Example 2.4 Let's look at our first Lean example

```
example (H_A : A) (H_B : B) : (A ∧ B) := And.intro H_A H_B
```

Lean aims to resemble the language used in mathematics. For instance, when defining a function or expression, one can use keywords such as **theorem** or **def**. Here, I used **example**, which is handy for defining anonymous expressions for demonstration purposes. After that comes the statement to be proved:

```
(H_A : A) (H_B : B) : (A ∧ B)
```

Meaning given a proof of A and a proof of B we can form a proof of $(A \wedge B)$. The operator `:=` assigns a value (or return an expression) for the statement which "has to be a proof of it". `And.intro` is implemented as:

And.intro: $p \rightarrow q \rightarrow (p \wedge q)$.

It says: if you give me a proof of p and a proof of q , then i return a proof of $p \wedge q$. We therefore conclude the proof by directly giving *And.intro* H_A H_B . Here another way of writing the same statment.

example $(H_p : p) (H_B : B) : \text{And}(A, B) := \langle H_p, H_B \rangle$

This system of inference rules allows us to construct proofs in an algorithmic and systematical way, organized in what is called a **proof tree**. To reduce complexity, we follow a **top-down** approach (see [Tho99] and [NPS90]). This methodology forms the basis of **proof assistants** like Lean, Coq, and Agda, which help verify the correctness of mathematical proofs by checking each step against these rules. We will see later that Lean, in fact, provides an info view of the proof tree which helps us understand and visualize the proof structure. Let's examine a concrete example of a proof.

Example 2.5 (Associativity of Conjunction) *We prove that $(A \wedge B) \wedge C$ implies $A \wedge (B \wedge C)$. First, from the assumption $(A \wedge B) \wedge C$, we can derive A :*

$$\frac{(A \wedge B) \wedge C}{\frac{A \wedge B}{A} \wedge E_1} \wedge E_1$$

Second, we can derive $B \wedge C$:

$$\frac{\frac{(A \wedge B) \wedge C}{\frac{A \wedge B}{B} \wedge E_2} \wedge E_1 \quad \frac{(A \wedge B) \wedge C}{C} \wedge E_2}{B \wedge C} \wedge I$$

Finally, combining these derivations we obtain $A \wedge (B \wedge C)$:

$$\frac{(A \wedge B) \wedge C \vdash A \quad (A \wedge B) \wedge C \vdash B \wedge C}{A \wedge (B \wedge C)} \wedge I$$

Example 2.6 (Lean Implementation) *Let us now implement the same proof in Lean.*

```
theorem and_associative : (A ∧ B) ∧ C → A ∧ (B ∧ C) :=
  fun h : (A ∧ B) ∧ C =>
    -- First, from the assumption (A ∧ B) ∧ C, we can derive A:
    have ab : A ∧ B := h.left -- extracts (derive) a proof of (A ∧ B)
    from the assumption
    have a : A := ab.left -- extracts A from (A ∧ B)
    -- Second, we can derive B ∧ C (here we only extract b and c and
    combine them in the next step)
    have c : C := h.right
    have b : B := ab.right
    -- Finally, combining these derivations we obtain A ∧ (B ∧ C)
    show A ∧ (B ∧ C) from ⟨a, ⟨b, c⟩⟩
```

Listing 1: Associativity of Conjunction in Lean

We introduce the `theorem` with the name `and_associative`, which can be referenced in subsequent proofs. The type signature $(A \wedge B) \wedge C \rightarrow A \wedge (B \wedge C)$ represents our logical implication. The `:=` operator introduces the proof term that establishes the theorem's validity. In the previous code example this proof was directly given. Here, we construct it using a function with the `fun` keyword. Why a function? We have already encountered the Curry-Howard correspondence in Lean previously, though without explicitly stating it. According to this correspondence, a proof of an implication can be understood as a function that takes a hypothesis as input and produces the desired conclusion as output. We will revisit this concept in more detail later. The `have` keyword introduces local lemmas within our proof scope, allowing us to break down complex reasoning into manageable intermediate steps, mirroring our natural deduction proof of before. Finally, the `show` keyword presents our final result.

To capture more complex mathematical ideas, we extend our system from propositional logic to **predicate logic**. A **predicate** is a statement or proposition that depends on a variable. In propositional logic we represent a proposition simply by P . In predicate logic, this is generalized: a predicate is written as $P(a)$, where a is a variable. Notice that a predicate is just a function. This extension allows us to introduce **quantifiers**: \forall ("for all") and \exists ("there exists"). These quantifiers express that a given formula holds either for every object or for at least one object, respectively. In Lean if α is any type, we can represent a predicate P on α as an object of type $\alpha \rightarrow \text{Prop}$. `Prop` stands for proposition, and it is an essential component of Lean's type system. For now, we can think of it as a special type whose inhabitants are proofs; somewhat paradoxically, a type of types. `Prop` stands for *proposition*, and it is an essential component of Lean's type system. Thus given an $x : \alpha$ (an element with type α) $P(x) : \text{Prop}$ would be representative of a proposition holding for x .

When introducing variables into a formal language we must keep in mind that the specific choice of a variable name can be substituted without changing the meaning of the predicate or statement. This should feel familiar from mathematics, where the meaning of an expression does not depend on the names we assign to variables. Some variables are **bound** (constrained), while others remain **free** (arbitrary, in programming often called "dummy" variables). When substituting variables, it is important to ensure that this distinction is preserved. This phenomenon, called **variable capture**, parallels familiar mathematical practice: if $f(x) := \int_1^2 \frac{dt}{x-t}$, then $f(t)$ equals $\int_1^2 \frac{ds}{t-s}$, not the ill-defined $\int_1^2 \frac{dt}{t-t}$. The same principle applies to predicate logic. For example, consider

$$\exists y. (y > x).$$

This states that for a given x there exists a y such that $y > x$. If we naively substitute $y + 1$ for x , we would obtain

$$\exists y. (y > y + 1),$$

where the y in $y + 1$ has been **captured** by the quantifier $\exists y$. This transforms the original statement from "there exists some y greater than the free variable

x " into the always-false statement "there exists some y greater than itself plus one." To avoid the problem, in the above example, we would first rename the bound variable to something fresh say z , obtaining $\exists z. (z > x)$, and then safely substitute to get $\exists z. (z > y + 1)$.

Notation 2.7 We use the notation $\phi[t/x]$ for *substitution*, meaning all occurrences of the free variable x in formula (or expression) ϕ are replaced by term t .

We can now present the inference rules for quantifiers.

Universal Quantification (\forall)

- Introduction

$$\frac{A}{\forall x. A} \forall I$$

The variable x must be arbitrary in the derivation of A . This rule captures statements like $\forall x \in \mathbb{N}, x$ has a successor, but would not apply to $\forall x \in \mathbb{N}, x$ is prime (since we cannot derive this for an arbitrary natural number).

- Elimination

$$\frac{\forall x. A}{A[t/x]} \forall E$$

The conclusion $A[t/x]$ represents the substitution of term t for variable x in formula A . From a proof of $\forall x. A(x)$ we can infer $A(t)$ for any term t .

Existential Quantification (\exists)

- Introduction

$$\frac{A[t/x]}{\exists x. A} \exists I$$

The substitution premise means that if we can find a specific term t for which $A(t)$ holds, then we can introduce the existential quantifier. The introduction rule requires a witness t for which the predicate holds.

- Elimination

$$\frac{\exists x. A \quad [A] \vdash B}{B} \exists E$$

To eliminate an existential quantifier, we assume A holds for some witness and derive B without making any assumptions about the specific witness.

We can give an informal reading of the quantifiers as infinite logical operations:

$$\begin{aligned}\forall x. A(x) &\equiv A(a) \wedge A(b) \wedge A(c) \wedge \dots \\ \exists x. A(x) &\equiv A(a) \vee A(b) \vee A(c) \vee \dots\end{aligned}$$

The expression $\forall x. P(x)$ can be understood as generalized form of implication. If P is any proposition, then $\forall x. P$ expresses that P holds regardless of the choice of x . When P is a predicate, depending on x , this captures the idea that we can derive P from any assumption about x . Moreover, there is a duality between universal and existential quantification. We shall develop all this discussions further after exploring their computational (type theoretical) meaning.

Example 2.8 *Lean expresses quantifiers as follow.*

```
∀ (x : X), P x
forall (x : X), P x -- another notation
```

Listing 2: For All

```
∃ (x : X), P x
exist (x : X), P x -- another notation
```

Listing 3: Exists

Where x is a variable with a type X , and $P\ x$ is a proposition, or predicate, holding for x .

Example 2.9 (Existential introduction in Lean) *When introducing an **existential** proof, we need a **pair** consisting of a witness and a proof that this witness satisfies the statement.*

```
example (x : Nat) (h : x > 0) : ∃ y, y < x :=
  Exists.intro 0 h -- or shortly ⟨0, h⟩
```

The **existential elimination rule** (`Exists.elim`) performs the opposite operation. It allows us to prove a proposition Q from $\exists x, P(x)$ by showing that Q follows from $P(w)$ for an **arbitrary** value w .

Example 2.10 (Existential elimination in Lean) *The existential rules can be interpreted as an infinite disjunction, so that existential elimination naturally corresponds to a **proof by cases** (with only one single case). In Lean, this reasoning is carried out using **pattern matching**, a known mechanism in functional programming for dealing with cases, with **let** or **match**, as well as by using **cases** or **rcases** construct.*

```
example (h : ∃ n : Nat, n > 0) : ∃ n : Nat, n > 0 :=
  match h with
  | ⟨witness, proof⟩ => ⟨witness, proof⟩
```

Example 2.11 The *universal quantifier* may be regarded as a generalized function. Accordingly, In Lean, universal elimination is simply function application.

```
example : ∀ n : Nat, n ≥ 0 :=
  fun n => Nat.zero_le n
```

Functions are primitive objects in type theory. For example, it is interesting to note that a relation can be expressed as a function: $R : \alpha \rightarrow \alpha \rightarrow \mathbf{Prop}$. Similarly, when defining a predicate ($P : \alpha \rightarrow \mathbf{Prop}$) we must first declare $\alpha : \mathbf{Type}$ to be some arbitrary type. This is what is called **polymorphism**. A canonical example is the identity function, written as $\alpha \rightarrow \alpha$, where α is a type variable. It has the same type for both its domain and codomain, this means it can be applied to booleans (returning a boolean), numbers (returning a number), functions (returning a function), and so on. In the same spirit, we can define a transitivity property of a relation as follows:

```
def Transitive (α : Type) (R : α → α → Prop) : Prop :=
  ∀ x y z, R x y → R y z → R x z
```

To use `Transitive`, we must provide both the type α and the relation itself. For example, here is a proof of transitivity for the less-than relation on \mathbb{N} (in Lean `Nat` or \mathbb{N}):

```
theorem le_trans_proof : Transitive Nat (· ≤ · : Nat → Nat → Prop) :=
  fun x y z h1 h2 => Nat.le_trans h1 h2 -- this lemma is provided by Lean
```

Looking at this code, we immediately notice that explicitly passing the type argument `Nat` is somewhat repetitive. Lean allows us to omit it by letting the type inference mechanism fill it in automatically. This is achieved by using **implicit arguments** with curly brackets:

```
def Transitive {α : Type} (R : α → α → Prop) : Prop :=
  ∀ x y z, R x y → R y z → R x z
theorem le_trans_proof : Transitive (· ≤ · : Nat → Nat → Prop) :=
  fun x y z h1 h2 => Nat.le_trans h1 h2
```

Lean's type inference system is quite powerful: in many cases, types can be completely inferred without explicit annotations. For instance, in earlier examples, Lean automatically inferred that the types of A , B , and C were `Prop`. Let us now revisit the transitivity proof, but this time for the less-than-equal relation on the rational numbers (`Rat` or \mathbb{Q}) instead.

```
import Mathlib

theorem rat_le_trans : Transitive (· ≤ · : Rat → Rat → Prop) :=
  fun _ _ h1 h2 => Rat.le_trans h1 h2
```

Here, `Rat` denotes the rational numbers in Lean, and `Rat.le_trans` is the transitivity lemma for \leq on rational numbers, provided by `Mathlib`. We import `Mathlib` to access `Rat` and `le_trans`. `Mathlib` is the community-driven mathematical library for Lean, containing a large body of formalized mathematics and

ongoing development. It is the de facto standard library for both programming and proving in Lean [Com20], we will dig into it as we go along. Notice that we used a function to discharge the universal quantifiers required by transitivity. The underscores indicate unnamed variables that we do not use later. If we had named them, say $x\ y\ z$, then: `h1` would be a proof of $x \leq y$, `h2` would be a proof of $y \leq z$, and `Rat.le_trans h1 h2` produces a proof of $x \leq z$. The `Transitive` definition is imported from `Mathlib` and similarly defined as before.

Example 2.12 *The code can be made more readable using tactic mode. This mode comes with an info view showing the goal to solve and proof structure. In this mode, you use tactics, commands provided by Lean or defined by users, to carry out proof steps succinctly, avoid code repetition, and automate common patterns. This often yields shorter, clearer proofs than writing the full term by hand.*

```
import Mathlib

theorem rat_le_trans : Transitive (· ≤ · : Rat → Rat → Prop) := by
  intro x y z hxy hyz
  exact le_trans hxy hyz
```

This proof performs the same steps but is much easier to read. Using `by` we enter Lean’s tactic mode, which (with the info view) shows the current goal and context. Move your cursor just before `by` and observe the info view change. The goal is shown displayed $\vdash \text{Transitive } \text{fun } x1\ x2 \mapsto x1 \leq x2$ at first. The tactic `intro` introduces the variables and hypotheses corresponding to the universal quantifiers and assumptions. Now position your cursor just before `exact` and observe the info view again. The goal is now $\vdash x \leq z$, with the context showing the variables and hypothesis introduced by the previous tactic. `exact` closes the goal by supplying the term `Rat.le_trans hxy hyz` that matches with the goal (the specification of `Transitive`). You can hover over each tactic to see its definition and documentation.

In these examples we have used predefined lemmas such as `Nat.le_trans` and `Rat.le_trans`, just to simplify the presentation. We can now dig in to the implementation of these lemmas. Let’s look at the source code of `Rat.le_trans`. The `Mathlib 4` documentation website is at: https://leanprover-community.github.io/mathlib4_docs, and The documentation for `Rat.le_trans` is at: https://leanprover-community.github.io/mathlib4_docs/Mathlib/Algebra/Order/Ring/Unbundled/Rat.html#Rat.le_trans Click the “source” link there to jump to the implementation in the `Mathlib` repository. In editors like VS Code you can also jump directly to the definition (Ctrl-click; Cmd-click on macOS). Lean has built-in types `Nat` (natural numbers), `Int` (integers), and `Rat` (rational numbers). While Lean provides basic arithmetic and order relations for these types, many advanced properties and theorems live in `Mathlib`. The proof uses several tactics and lemmas from `Mathlib`. You can follow the proof step by step using the info view in tactic mode, by positioning the cursor on each line and observing the changes in the goal and context. The `rw`

or `rewrite` tactic is very common and syntactically similar to the mathematical practice of rewriting an expression using an equality. In this case, with `at`, we use it to rewrite the hypotheses `hab` and `hbc` using the another Mathlib's lemma `Rat.le_iff_sub_nonneg`, which states that for any two rational numbers `x` and `y`, $x \leq y$ is equivalent to $0 \leq y - x$. Thus we now have the hypotheses tranformed to :

```
hab : 0 ≤ b - a
hbc : 0 ≤ c - b
```

The `have` tactic introduces an intermediate result. If you omit a name, Lean assigns it the default name `this`. In our situation, from `hab : a ≤ b` and `hbc : b ≤ c` we can derive that `b - a` and `c - b` are nonnegative, hence their sum is nonnegative:

```
this : 0 ≤ b - a + (c - b)
```

The most involved step uses `simp_rw` to simplify the expression via a sequence of other existing Mathlib's lemmas. The tactic `simp_rw` is a variant of `simp`: it performs rewriting using the `simp` set (and any lemmas you provide), applying the rules in order and in the given direction. Lemmas that `simp` can use are typically marked with the `@[simp]` attribute. This is particularly useful for simplifying algebraic expressions and equations. After these simplifications we obtain:

```
this : 0 ≤ c - a
```

Clearly, the proof relies mostly on `Rat.add_nonneg`. Its source code is fairly involved and uses advanced features that are beyond our current scope. Nevertheless, it highlights an important aspect of formal mathematics in Mathlib. Mathlib defines `Rat` as an instance of a linear ordered field, implemented via a normalized fraction representation: a pair of integers (numerator and denominator) with positive denominator and coprime numerator and denominator [Lea25b]. To achieve this, it uses a **structure**. In Lean, a structure is a dependent record type used to group together related fields or properties as a single data type. Unlike ordinary records, the type of later fields may depend on the values of earlier ones. Defining a structure automatically introduces a constructor (usually `mk`) and projection functions that retrieve (deconstruct) the values of its fields. Structures may also include proofs expressing properties that the fields must satisfy.

In order to work with rational numbers in Mathlib, we use the `Rat.mk`' constructor to create a rational number from its numerator and denominator, if omitted the default would be `Rat.mk`. The fields `den_nz` and `reduced` are proofs that the denominator is nonzero and that the numerator and denominator are coprime, respectively. These proofs are automatically generated by Lean's `decide` tactic, which can solve certain decidable propositions (to be discussed in the next section).

Example 2.13 *Here is how we can define and manipulate rational numbers in Lean.*

When working with rational numbers, or more generally with structures, we must provide the required proofs as arguments to the constructor (or Lean must be able to ensure them). For instance `Rat.mk' 1 0` or `Rat.mk' 2 6` would be rejected. In the case of rationals, Mathlib unfolds the definition through `Rat.numDenCasesOn`. This principle states that, to prove a property of an arbitrary rational number, it suffices to consider numbers of the form n / d in canonical (normalized) form, with $d > 0$ and $\gcd n d = 1$. This reduction allows mathlib to transform proofs about \mathbb{Q} into proofs about \mathbb{Z} and \mathbb{N} , and then lift the result back to rationals.

Example 2.14 *I will present a simplified version of this implementation.*

In this version, we open the `Rat` namespace to access its definitions and lemmas directly (Notice that i use `num_nonneg` instead of `Rat.num_nonneg` in the next line). The proof begins by introducing the hypotheses `ha` and `hb` that `a` and `b` are nonnegative. The `rw ... at` tactic rewrites these hypotheses using the lemma `num_nonneg`, which states that a rational number is nonnegative if and only if its numerator is nonnegative. We use \leftarrow to indicate the direction of rewriting (from right to left). Next, we express the rational numbers in terms of their numerator and denominator using `num_divInt_den`, and apply the addition formula for rational numbers represented as `divInt_add_divInt`. The goal is then to prove that the resulting numerator is nonnegative. `num_divInt_den` transforms a rational number `r` into the form $r.num / \uparrow r.den$. The \uparrow symbol denotes the coercion from natural numbers to integers (remember in our definition of `Rat`, the numerator is an integer and the denominator a natural number, but here we need to translate everything to integers). We now have 3 goals. We use `divInt_nonneg_iff_of_pos_right` to reduce this to proving that the numerator is nonnegative, given that the denominator is positive. This requires two subgoals: proving the numerator is nonnegative and the denominator is positive. For the numerator, we use `Int.add_nonneg` to show that the sum of two nonnegative integers is nonnegative. For the denominator, we first translate the problem from integers to natural numbers, using `norm_cast`. and use `Nat.mul_pos` to show that the product of two positive natural numbers is positive. Finally, we use `norm_cast` to handle the necessary type casts between integers and natural numbers automatically and close the remaining goals with the nonzero denominator conditions given from the `Rat` structure (`den_nz`). Lean encourages to separate subgoals with `·` and proper indentation, making the corresponding proof more readable.

We made extensive use of type casting and coercions in this proof, handled by the `norm_cast` tactic, which requires some explanation ([LM20]). Lean type system lack of subtypes means that types like `Nat`, `Int`, and `Rat` are distinct and do not have a subtype relationship. In order to translate between these types, we need to use explicit type casts or coercions. For example, natural numbers (`Nat`) can be coerced to integers (`Int`) and integers can be coerced to rational numbers (`Rat`). The `norm_cast` tactic simplifies expressions involving such coercions by normalizing them, making it easier to reason about mixed-type expressions. It will be otherwise a long and tedious process to manually insert and manage

these coercions throughout the proof. `norm_cast` is another example of a tactic that leverages Lean’s metaprogramming capabilities to automate common proof patterns. (I CAN DISCUSS THIS FURTHER IF NEEDED).

The theorem previously used with natural numbers, `Nat.le_trans`, is part of Lean’s internal library at `/lean/Init/Prelude.lean`. Mathlib is built on top of this base library. More generally, the transitivity property holds not only for naturals but also for integers, reals, and, in fact, for any partially ordered set. Mathlib provides a general lemma `le_trans` for any type α endowed with partial ordering. This is achieved through type classes, Lean’s mechanism for defining and working with abstract algebraic structures in an ad hoc polymorphic manner. Type classes provide a powerful and flexible way to specify properties and operations that can be shared across different types, thereby enabling polymorphism and code reuse. Ad hoc polymorphism arises when a function is defined over several distinct types, with behavior that varies depending on the type. A standard example ([WB89]) is overloaded multiplication: the same symbol denotes multiplication of integers (e.g. `3 * 3`) and of floating-point numbers (e.g. `3.14 * 3.14`). By contrast, parametric polymorphism occurs when a function is defined over a range of types but acts uniformly on each of them. For instance, the length function applies in the same way to a list of integers and to a list of floatingpoints.

Under the hood, a type class is a structure. An important aspect of structures, and hence type classes, is that they are powered by hierarchy and composition. For example, a monoid is a semigroup with an identity element, and a group is a monoid with inverses. In Lean, we can express this by defining a `Monoid` structure that extends the `Semigroup` structure, and a `Group` structure that extends the `Monoid` structure using the `extends` keyword.

The symbol `*` on $(\alpha : \text{Type}^*)$ indicates a universe variable (we will discuss universes later). Sometimes, in order to avoid inconsistencies between types (like Girard’s paradox), universes must be specified explicitly. This is an example of universe polymorphism, thus we have seen all the polymorphism flavors in Lean. On the other hand, Type classes are defined using the `class` keyword, which is syntactic sugar for defining a structure. Thus the previous example can be rewritten similarly, using type classes: The main difference is that type classes support **instance resolution**, using the keyword `instance` to declare that a particular type is an instance of a type class, which inherits the properties and operations defined in the type class.

Let’s look at a more concrete example, say we would like to define a generic `Transitive` structure with field `le_trans`:

```
structure TransitiveProperty ( $\alpha : \text{Type}$ ) where
  le_trans :  $\alpha \rightarrow \alpha \rightarrow \alpha$ 
```

Now in overload different types such as `Rat`, `INat` and `Nat` we would define:

This idea is extensively used in Mathlib to define and work with algebraic structures.

Instances of a type class can be automatically inferred by Lean’s type inference system, allowing for concise and expressive code. This mechanism is

particularly useful for defining and working with algebraic structures, such as groups, rings, and fields, as well as order structures like preorders and partial orders. Mathematically, a partially ordered set consists of a set P and a binary relation \leq on P that is transitive and reflexive ([Lea25a] Structures)

The `class Preorder` declares a type class over a type α , bundling the \leq and $<$ relations (inherited via `extends LE alpha, LT alpha`) with the pre-order axioms: reflexivity (`le_refl`) and transitivity (`le_trans`). The theorem `lt_iff_le_not_ge` provides a characterization of the strict order, proved automatically (by `intros; rfl`). The `instance` declaration connects the `Preorder` class to Lean’s `Grind` tactic automation, which allows automatic reasoning with preorder properties.

This design pattern is the foundation of Lean’s powerful mathematical library, allowing complex abstract algebraic and order structures to be expressed succinctly and compositionally.

References

- [Com20] The mathlib Community. “The Lean Mathematical Library”. In: *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs (CPP ’20)*. New Orleans, LA, USA: ACM, 2020, pp. 367–381. DOI: 10.1145/3372885.3373824. URL: <https://leanprover-community.github.io/papers/mathlib-paper.pdf>.
- [GTL89] Jean-Yves Girard, Paul Taylor, and Yves Lafont. *Proofs and Types*. Vol. 7. Cambridge Tracts in Theoretical Computer Science. Translated and adapted from the French book “Le Point Aveugle”. Cambridge University Press, 1989.
- [Lea25a] Lean Mathematical Library community. *Mathematics in Lean*. https://leanprover-community.github.io/mathematics_in_lean/mathematics_in_lean.pdf. Version v4.23.0-rc2. 2025. URL: https://leanprover-community.github.io/mathematics_in_lean/mathematics_in_lean.pdf.
- [Lea25b] Lean Mathematical Library community. *mathlib — The Lean Mathematical Library*. <https://github.com/leanprover-community/mathlib>. Version v4.23.0-rc2. 2025. URL: <https://github.com/leanprover-community/mathlib>.
- [LM20] Robert Y. Lewis and Paul-Nicolas Madelaine. “Simplifying Casts and Coercions”. In: *arXiv preprint arXiv:2001.10594v2* (2020). arXiv: 2001.10594 [cs.PL]. URL: <https://arxiv.org/abs/2001.10594>.
- [NPS90] Bengt Nordström, Kent Petersson, and Jan M. Smith. *Programming in Martin-Löf’s Type Theory: An Introduction*. Now available from <https://www.cse.chalmers.se/research/group/logic/book/book.pdf>. Oxford, UK: Oxford University Press, 1990. URL: <https://www.cse.chalmers.se/research/group/logic/book/book.pdf>.

- [Tho99] Simon Thompson. *Type Theory & Functional Programming*. March 1999. Computing Laboratory, University of Kent, 1999.
- [Wad15] Philip Wadler. “Propositions as Types”. In: *Communications of the ACM* 58.12 (2015), pp. 75–84. DOI: 10.1145/2699407. URL: <https://homepages.inf.ed.ac.uk/wadler/papers/propositions-as-types/propositions-as-types.pdf>.
- [WB89] Philip Wadler and Stephen Blott. “How to Make Ad-Hoc Polymorphism Less Ad Hoc”. In: *16th ACM Symposium on Principles of Programming Languages (POPL)*. Austin, TX, USA: ACM, Jan. 1989, pp. 1–12. URL: <https://web.engr.oregonstate.edu/~walkiner/teaching/cs583-sp21/files/Wadler-TypeClasses.pdf>.